



Firmware Security

Roland Nagy
CrySyS Lab, HIT, BME
rnagy@crysys.hu

Agenda

- Intro
- Secure Boot
 - Goals & requirements
 - Crypto recap
 - Authentication schemes & hardware support
 - Example 1: boot flow of a microcontroller
 - Example 2: boot flow of UEFI
- Secure Firmware Update
 - Goals & requirements
 - A/B partitioning (Seamless Update) & rollbacks
 - Example 1: secure update of a microcontroller
 - Example 2: secure update of UEFI

Introduction

- What is the firmware?
 - “software for hardware”
 - “software that provides low level control for hardware”
 - “software that runs before the OS”
- For desktops/servers -> BIOS/UEFI
 - Starts a bootloader/OS
- For embedded devices, we often don't have an OS
 - We call firmware everything that is code
- For network devices, again, we often call everything firmware
 - Despite having an OS, which in this case is part of the firmware image

Introduction

- So what is the firmware?
 - Code that runs when the device is starting
 - Perform hardware initialization
 - Starts other software components
- Why firmware security is important?
 - OS security features not available yet
 - If compromised, other components started by the firmware cannot be considered safe and secure either
 - After all, it's just code, it can have vulnerabilities
 - Modern firmwares can be quite large
 - Malware target it
 - » Mebromi for BIOS
 - » BlackLotus for UEFI
 - » Boot sector viruses since 1986 (Brain, Lehigh, SCA, Ping-Pong virus, etc.)

Secure Boot

Secure Boot

- Very, very high level boot flow:



Secure Boot - Authentication

- To authenticate the loaded software component, we can use authentication schemes:
 - Hash-based authentication
 - » Using hash functions
 - MAC-based authentication
 - » Using Message Authentication Codes
 - Signature-based authentication
 - » Using digital signatures

Cryptographic hash functions

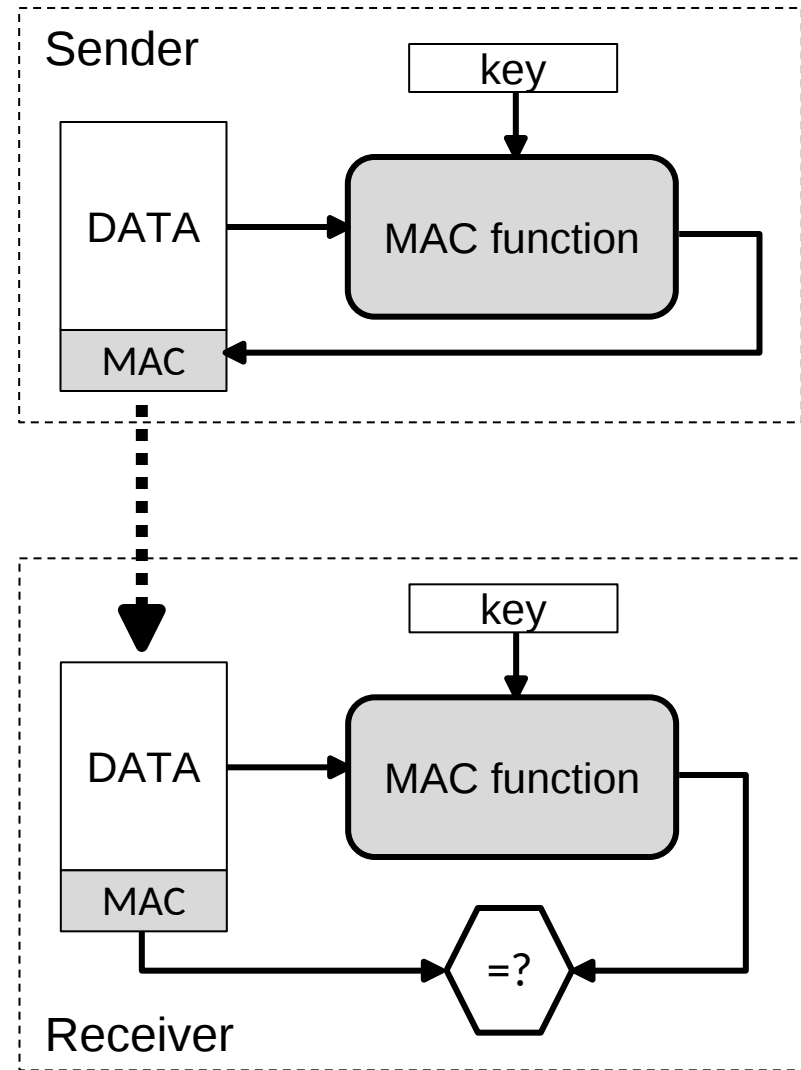
- a hash function is a function that maps arbitrary long messages into a fixed length output (n bits)
- notation and terminology:
 - x – (input) message
 - $y = H(x)$ – hash value, message digest, fingerprint
- typical applications:
 - the hash value of a message can serve as a compact representative image of the message (similar to fingerprints)
 - increase the efficiency of digital signatures by signing the hash instead of the message (expensive operation is performed on small data)
 - password hashing
 - **checking if data was modified**
- examples:
 - (~~MD5, SHA-1~~) SHA-2 and SHA-3

Properties of crypto hash functions

- ease of computation
 - given an input x , the hash value $H(x)$ of x is easy to compute
- **weak collision resistance** (2^{nd} preimage resistance)
 - given an input x , it is computationally infeasible to find a second input x' such that $H(x') = H(x)$
- **strong collision resistance** (collision resistance)
 - it is computationally infeasible to find any two distinct inputs x and x' such that $H(x) = H(x')$
- **one-way property** (preimage resistance)
 - given a hash value y (for which no preimage is known), it is computationally infeasible to find any input x such that $H(x) = y$
- collision resistant hash functions can typically be modeled as a random function (similar to block ciphers)

Message Authentication Codes (MAC)

- a MAC function is a function that maps an arbitrary long message and a key (k bits) into a fixed length output (n bits)
 - can be viewed as a hash function with an additional input (the key)
- services:
 - **message authentication and integrity protection:** after successful verification of the MAC value, the receiver is assured that the message has been generated by the sender and it has not been altered in transit
- examples:
 - HMAC, CBC-MAC, CMAC

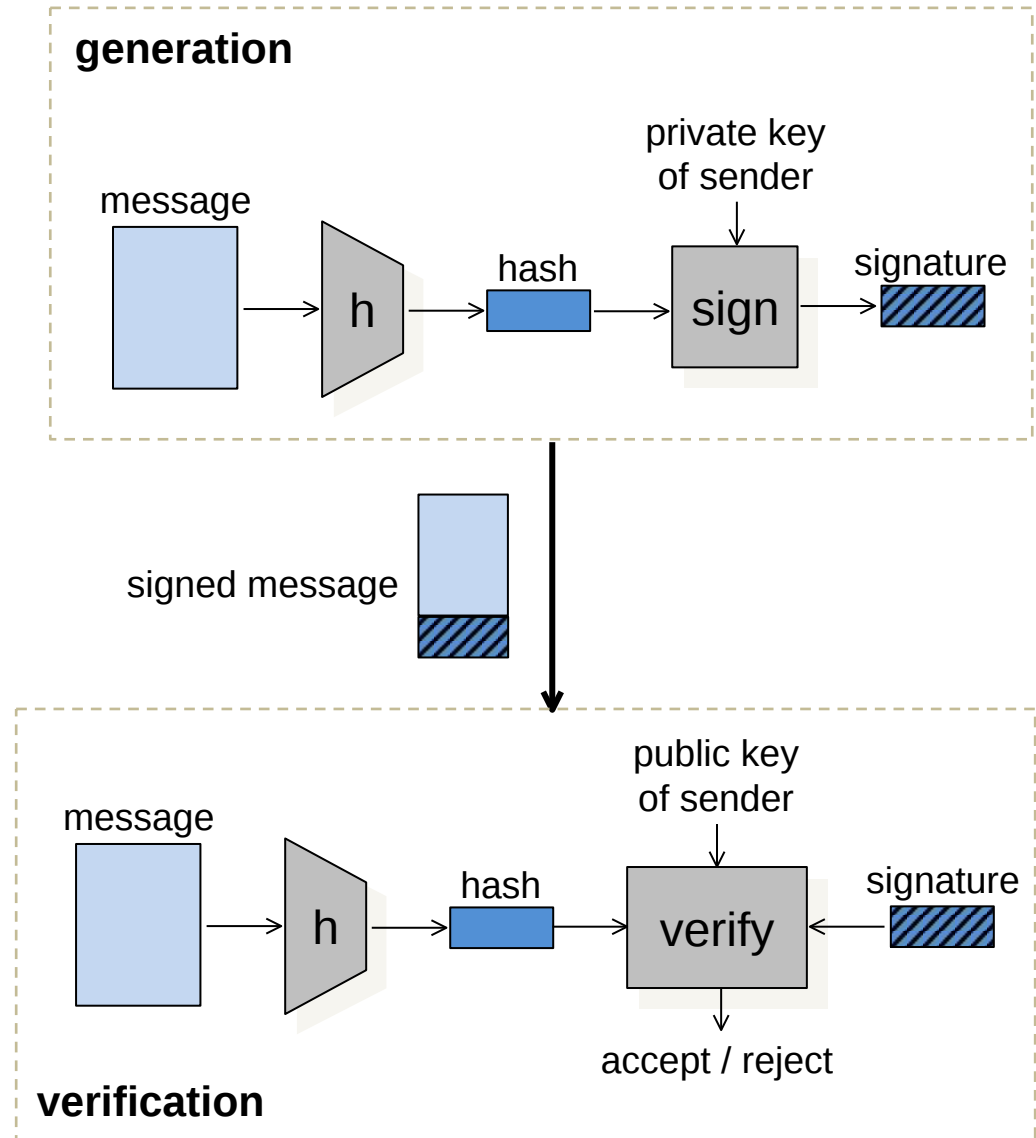


Digital signature schemes

- similar to MACs but they are
 - unforgeable by the receiver
 - verifiable by a third party
- services:
 - **message authentication and integrity protection:** after successful verification of the signature, the receiver is assured that the message has been generated by the sender and it has not been altered
 - **non-repudiation of origin:** the receiver can prove this to a third party (hence the sender cannot repudiate)
- examples: RSA, DSA, ECDSA


Hash-and-sign paradigm

- public/private key operations are slow
- increase efficiency by signing the hash of the message instead of the message
- it is essential that the hash function is collision resistant



Basic idea of certificates

- name and public key is linked together by the digital signature of a **trusted entity** called **Certification Authority (CA)**



- in order to verify a certificate you need to have an authentic copy of the public key of the CA
- advantage: only the CA's public key need to be distributed via out-of-band channels (scales better)

Certificates illustrated

Certificate Viewer: "Builtin Object Token: Network Solutions Certificate Authority"

General | Details

This certificate has been verified for the following uses:

SSL Certificate Authority

Issued To

Common Name (CN) Network Solutions Certificate Authority
Organization (O) Network Solutions L.L.C.
Organizational Unit (OU) <Not Part Of Certificate>
Serial Number 57:CB:33:6F:C2:5C:16:E6:47:16:17:E3:90:31:68:E0

Issued By

Common Name (CN) Network Solutions Certificate Authority
Organization (O) Network Solutions L.L.C.
Organizational Unit (OU) <Not Part Of Certificate>

Period of Validity

Begins On 2006. december 1.
Expires On 2030. január 1.

Fingerprints

SHA-256 Fingerprint 15:F0:BA:00:A3:AC:7A:F3:AC:88:4C:07:2B:10:11:A0:77:BD:77:C0:97:F4:01:64:B2:F8:59:8A:BD:83:86:0C
SHA1 Fingerprint 74:F8:A3:C3:EF:E7:B3:90:06:4B:83:90:3C:21:64:60:20:E5:DF:CE

Certificate Viewer: "Builtin Object Token: Network Solutions Certificate Authority"

General | **Details**

Certificate Hierarchy

Network Solutions Certificate Authority

Certificate Fields

- Subject
- Subject Public Key Info
 - Subject Public Key Algorithm
 - Subject's Public Key
- Extensions
 - Certificate Subject Key ID
 - Certificate Key Usage
 - Certificate Basic Constraints

Field Value

Modulus (2048 bits):

```
e4 bc 7e 92 30 6d c6 d8 8e 2b 0b bc 46 ce e0 27
96 de de f9 fa 12 d3 3c 33 73 b3 04 2f bc 71 8c
e5 9f b6 22 60 3e 5f 5d ce 09 ff 82 0c 1b 9a 51
50 1a 26 89 dd d5 61 5d 19 dc 12 0f 2d 0a a2 43
5d 17 d0 34 92 20 ea 73 cf 38 2c 06 26 09 7a 72
f7 fa 50 32 f8 c2 93 d3 69 a2 23 ce 41 b1 cc e4
d5 1f 36 d1 8a 3a f8 8c 63 e2 14 59 69 ed 0d d3
7f 6b e8 b8 03 e5 4f 6a e5 98 63 69 48 05 be 2e
```

Export...

Close

Certification Authority (CA)

- collection of hardware, software, and staff (people)
- main functions:
 - issues certificates for users or other CAs
 - maintains certificate revocation information
 - publishes currently valid certificates and certificate revocation lists (CRL)
 - maintains archives
- must comply with strict security requirements related to the protection and usage of its private keys (basis of trust)
 - uses tamper resistant Hardware Security Modules that enforce security policies (access and usage control)
 - defines and publishes its certificate issuing policies
 - complies with laws and regulations
 - is subject to regular control (by national supervising authority)

Secure Boot - Authentication

- To authenticate the loaded software component, we can use authentication schemes:
 - Hash-based authentication
 - » Using hash functions
 - » **Reference hash must be stored write protected**
 - MAC-based authentication
 - » Using Message Authentication Codes
 - » **Symmetric key must be stored write protected and confidential**
 - Signature-based authentication
 - » Using digital signatures
 - » **Signing public key, CA public key or their hash must be stored write protected**

Secure Boot – Hardware support

- To be able to use the authentication schemes, some hardware support is needed
 - One-time programmable memory (OTP)
 - » Once written, after that, it's read-only
 - Secure Hardware Extension (SHE)
 - » A cryptographic coprocessor
 - » Can perform block cipher operations
 - » Can store a symmetric key confidentially
 - Hardware Security Module
 - » Can be a crypto coprocessor or a separate device
 - » Can perform all sorts cryptographic operations
 - Hashing, symmetric/asymmetric encryption/decryption, MAC computation, etc.
 - » Can store any type of keys & hashes
 - » Usually has some sort of tamper protection

Secure Boot – Compatibility

- Different auth. schemes have different requirements, thus they are compatible with different hardware components

	Hash	MAC	Signature
OTP	Y*	N	Y**
SHE	N	Y	N
HSM	Y	Y	Y

*: Firmware image is not updatable

**: Image signing key or CA root certificate is not updatable


Secure Boot - RoT/CoT

- Root of Trust
 - A component we trust
 - Usually we are not checking its integrity
 - Trust is often provided by storing it in ROM

- Chain of Trust
 - The first component we trust, because it is in ROM
 - It checks the integrity of the next one, before control is passed to it
 - Thus we trust the second, because the first said it's fine & we trust it

 - For every component i , component $i-1$ must check its integrity
 - We trust i if we the check succeeded and we trust $i-1$

 - Thus we can build a Chain of Trust from the Root of Trust



Secure Boot on ESP32

Secure Boot on a microcontroller (ESP32)

- High-level boot flow
 - Boot ROM starts, checks if secure boot enabled
 - Boot ROM authenticates the bootloader
 - If the signature is valid, the bootloader is started
 - The bootloader authenticates the application image
 - If the signature is valid, the application is started
 - If any of the checks fail, the boot process is interrupted

Hardware support – ESP32

- ESP32 contains 4 eFuses (One-Time Programmable memory blocks)
 - 256 bit each, 8 * 32 bit blocks
 - BLK0: used entirely to store system configuration, like if secure boot or firmware encryption enabled
 - BLK1: used to store a symmetric key used by firmware encryption
 - » only accessible to HW
 - BLK2: used to store the SHA256 hash of the public key used by secure boot
 - BLK3: available for the application

Detailed Boot Flow 1 – ESP32

- On start-up, ROM code checks if the secure boot is enabled in BLK0
- If enabled, the Signature Block of the bootloader is checked
 - Magic byte & CRC checked
- If the Signature Block is valid, the bootloader image is checked
 - The hash of the key in the Signature Block is compared to the hash stored in BLK2
 - The image is hashed and the hash is compared to the one stored in the Signature Block
 - If the hash is correct, the signature is checked
- If all checks pass, the bootloader is loaded into memory & executed

Detailed Boot Flow 2 – ESP32

- The bootloader checks the Signature Block of the application image
 - Magic byte & CRC checked
- If the Signature Block is valid, the app image is checked
 - The hash of the key in the Signature Block is compared to the hash stored in BLK2
 - The image is hashed and the hash is compared to the one stored in the Signature Block
 - If the hash is correct, the signature is checked
- If all checks pass, the application is loaded into memory & executed
 - The bootloader may check multiple app images, until a valid one is found



UNIFIED EXTENSIBLE
FIRMWARE INTERFACE

UEFI

UEFI - History

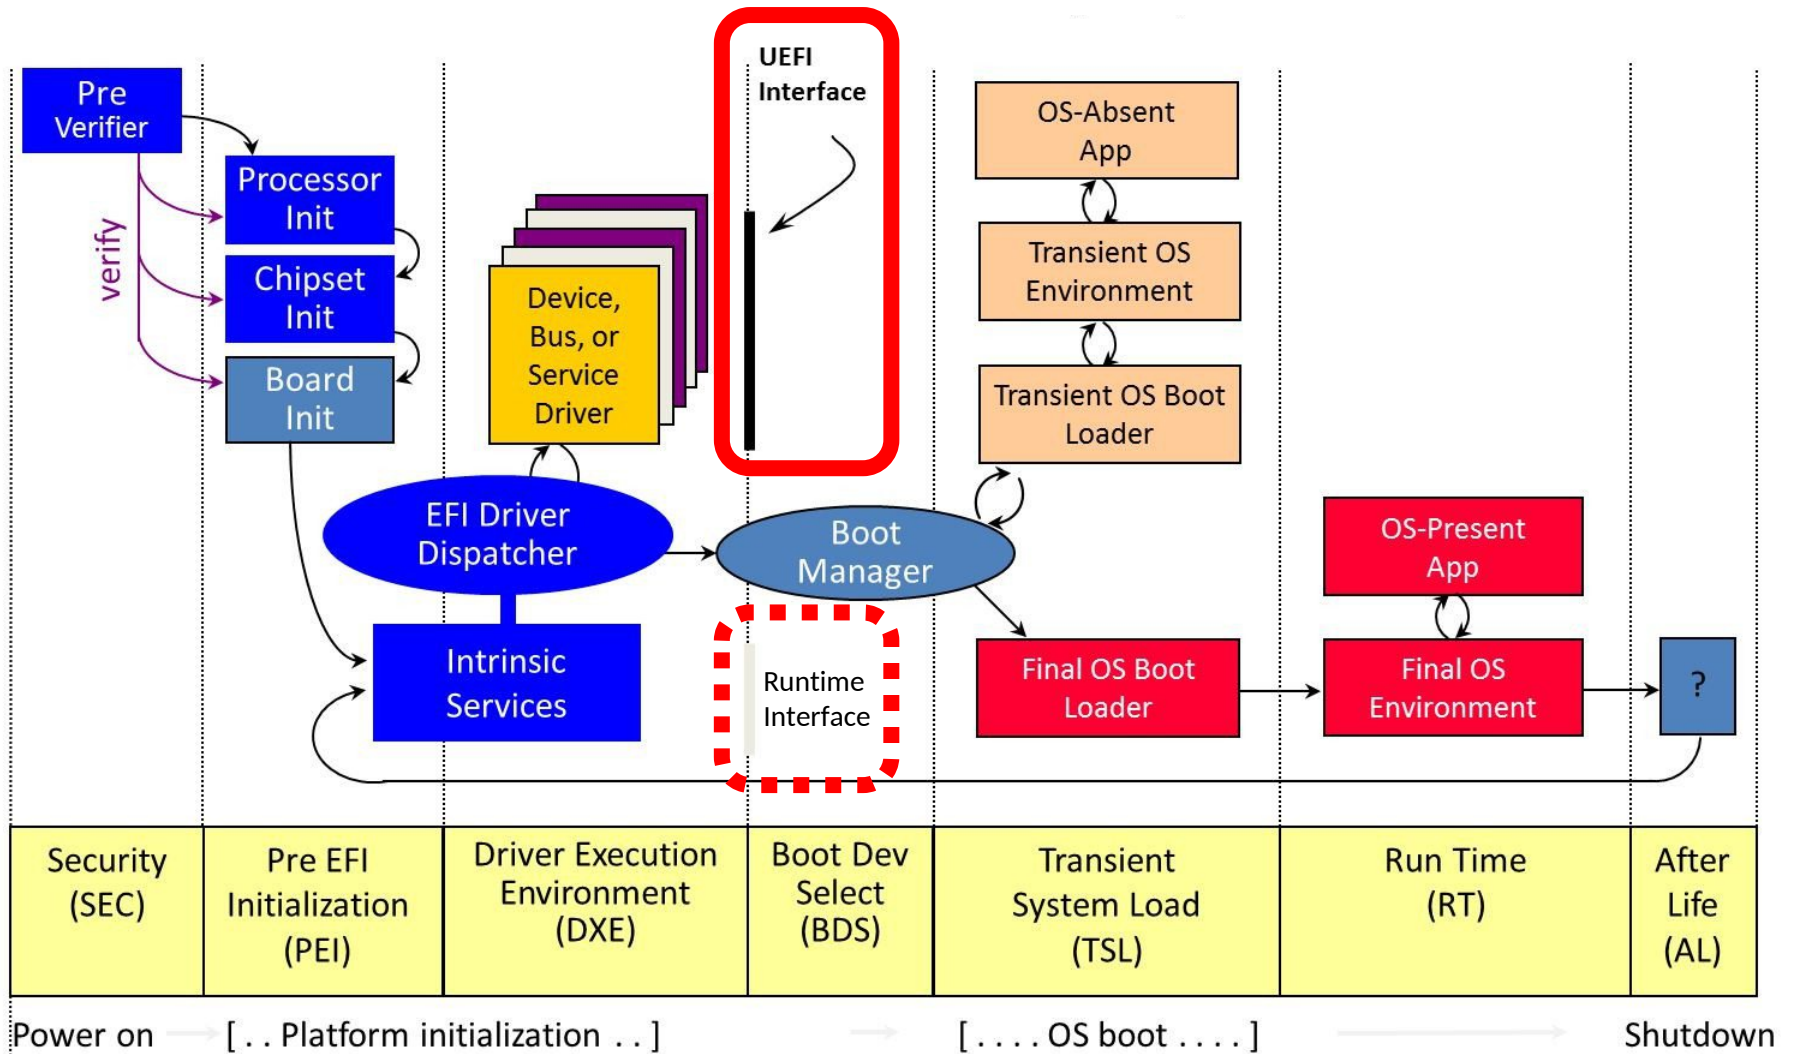
- A long time ago, in a galaxy far, far away...
 - ... there was BIOS (1975)
 - » Basic Input/Output System
 - » Developed by a handful of companies
 - Proprietary software
 - 0 compatibility between the different implementations
 - » Simple boot flow
 - Hardware check
 - Start whatever is at the first sector of the first disk
 - » Secure boot not supported
 - » Some serious limitations
 - Must run in 16-bit mode
 - 1 MB of memory for code
 - Cannot boot drives larger than 2 TB
 - » Intel, mid-1990s: „We need something better”



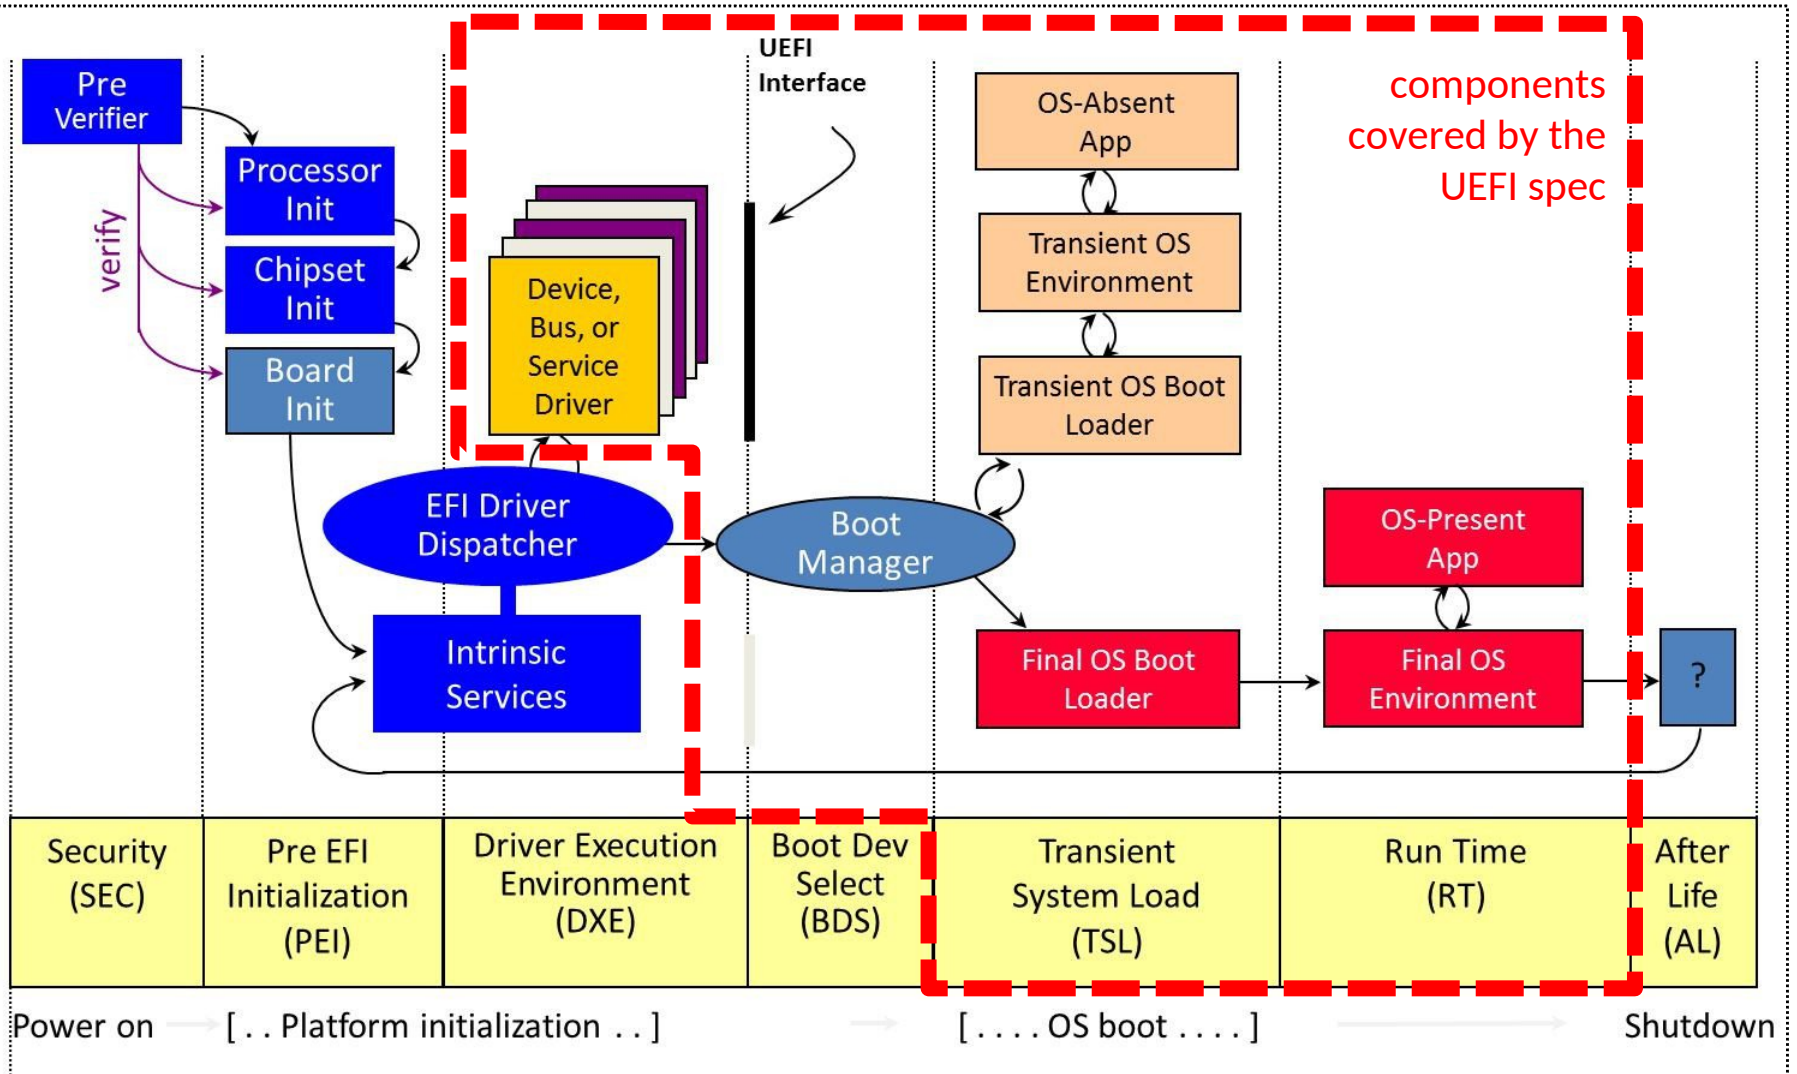
UEFI - History

- 1998 – Intel Boot Initiative
 - Later renamed to EFI
- 2004 – first open source UEFI implementation (Tiano by Intel)
- 2005 – EFI 1.10 were given to Unified EFI Forum
 - An alliance between tech companies to coordinate the specification of UEFI
 - » AMD, ARM, Apple, HP, Intel, Lenovo, Microsoft, etc.
 - Original EFI spec. owned by Intel, UEFI spec. owned by UEFI Forum
- 2006 Jan. – UEFI v2.0 with crypto & security features
- 2007 Jan. – UEFI v2.1 with network authentication & UI
- 2022 Aug. – Latest UEFI spec., v2.10

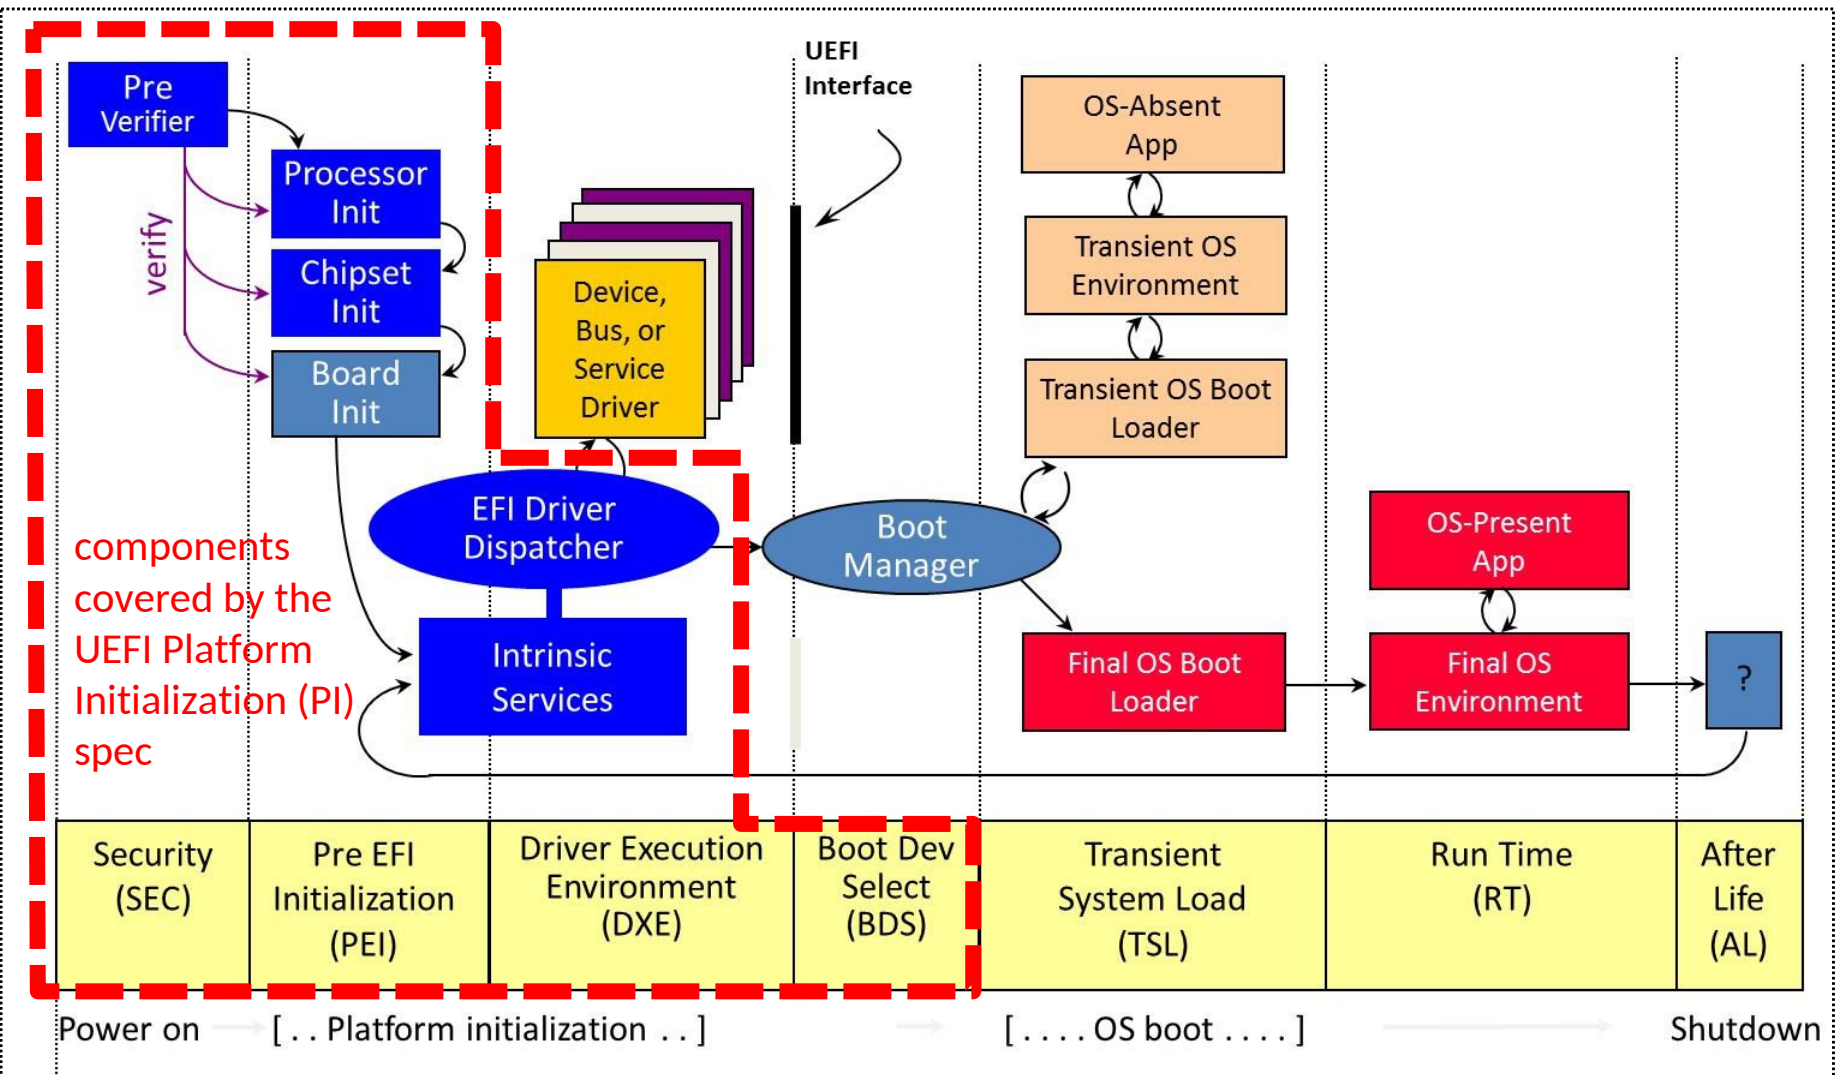
Platform boot flow overview



Platform boot flow overview

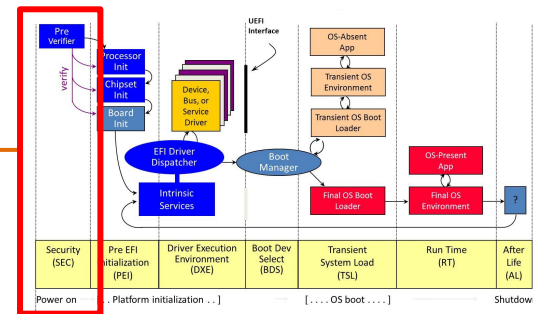


Platform boot flow overview



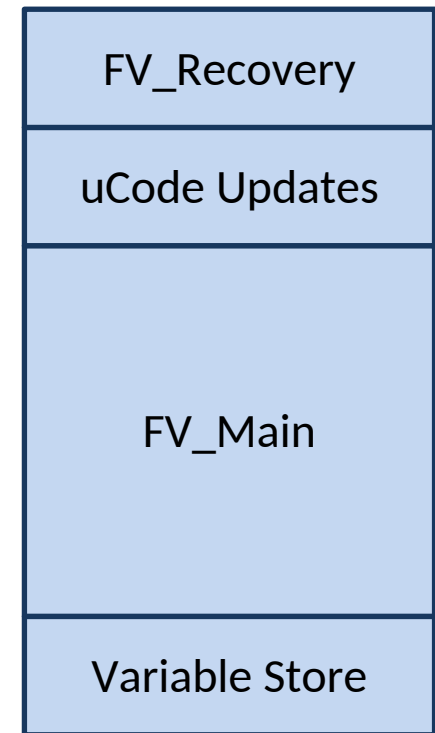
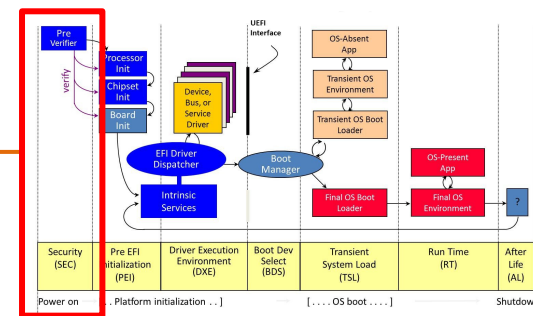
UEFI boot process – SEC

- Very first code executed by the CPU
 - Minimal code fetched directly from an SPI flash
 - Typically hand-coded assembly
 - Architecturally dependent and not portable
- SPI flash stuff
 - The CPU fetches the first instruction from address 0xFFFFFFF0 (re-directed to the flash by hardware)
 - This is just a JMP instruction to the start of the platform initialization code in the flash
- The SEC phase is responsible for
 - Handling all platform reset events (power-on, wake-up)
 - Executing microcode patch update to the CPU
 - Configuring the CPU Cache as RAM (CAR)
 - Note that the RAM has not yet been configured, but firmware code in later phases needs a C like execution environment (e.g., a stack to support function calls)
 - » we need some sort of RAM



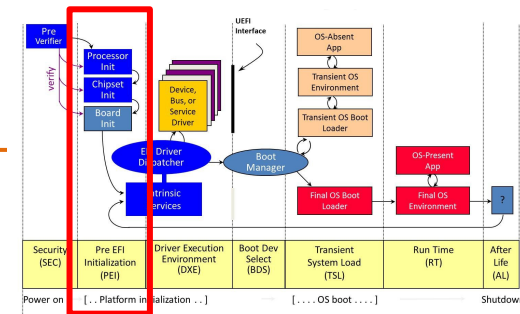
SPI flash content

- Firmware Volumes (FVs)
 - Logical firmware devices
 - Different boot phase code (e.g., PEI, DXE) may be stored on different volumes
- Each FV is organized into a Firmware File System (FFS)
 - A FFS contains files and their meta-data
 - Executable files can be in PE (Portable Executable) or TE (Terse Executable) format
- FV_Recovery
 - Contains the Boot Block, which holds the SEC and PEI phase code
- Other FVs contain compressed UEFI drivers (DXE phase) and remaining UEFI code



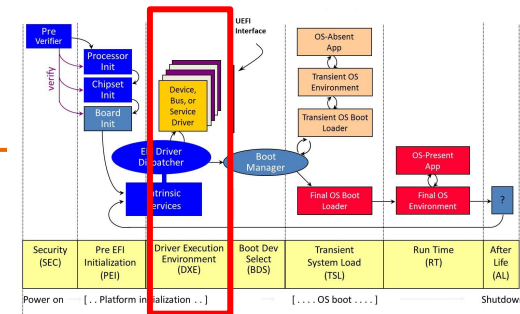
UEFI boot process – PEI

- Control is transferred to PEI phase code
 - PEI = Pre-EFI Initialization
 - Still fetched directly from the SPI flash
- PEI Dispatcher invokes PEI Modules (PEIMs) that perform early hardware and memory initialization (CPU, chipset, board init)
- Last PEIM called is DXE IPL (Initial Program Load), which decompresses FV_Main into main memory (RAM) and transitions to the DXE phase




UEFI boot process – DXE

- Control is transferred to DXE phase code
 - DXE = Driver Execution Environment
 - Code is executed from RAM
- DXE Dispatcher dispatches DXE and Runtime (RT) drivers from FV_Main (in RAM) into main memory (in RAM)
- The DXE phase is responsible for
 - Additional hardware initialization and configuration performed by DXE drivers
 - System Management Mode (SMM) setup
 - Secure Boot enforcement
 - Firmware update signature checks



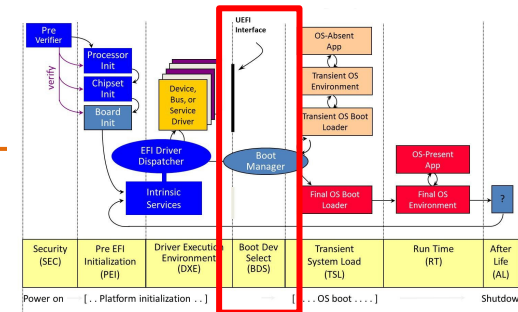
Setting up SMM and RT services



- SMM:
 - A platform-specific driver configures SMRAM and launches the SMM Dispatcher
 - The SMM Dispatcher loads SMM drivers from FV_Main into main memory and executes them
 - Some SMM drivers install SMI interrupt handlers
 - Typically performs tasks like power management and hardware control
 - Should only be used by the firmware
 - Transparent to the OS (Ring -2)
- RT services:
 - Some RT drivers install services callable by the OS at runtime
 - Note that DXE drivers are unloaded after OS boot, leaving only RT and SMM drivers to be available at runtime

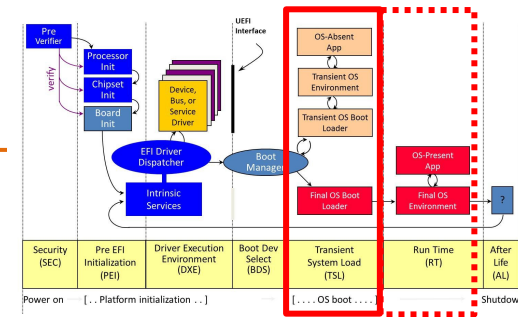
UEFI boot process – BDS

- Control is transferred to BDS phase code
 - BDS = Boot Device Selection
 - Typically, the BDS phase code is encapsulated in a single file loaded by the DXE phase
- The Boot Manager consults configuration information to decide where the OS should be booted from
- It has access via the UEFI interface to all UEFI Boot Services that the DXE phase set up --» it can use them to access the file system on the hard drive in order to find an OS bootloader
- If UEFI Secure Boot is turned on, it also checks the integrity of the OS bootloader before starting it



UEFI boot process – TSL (and RT)

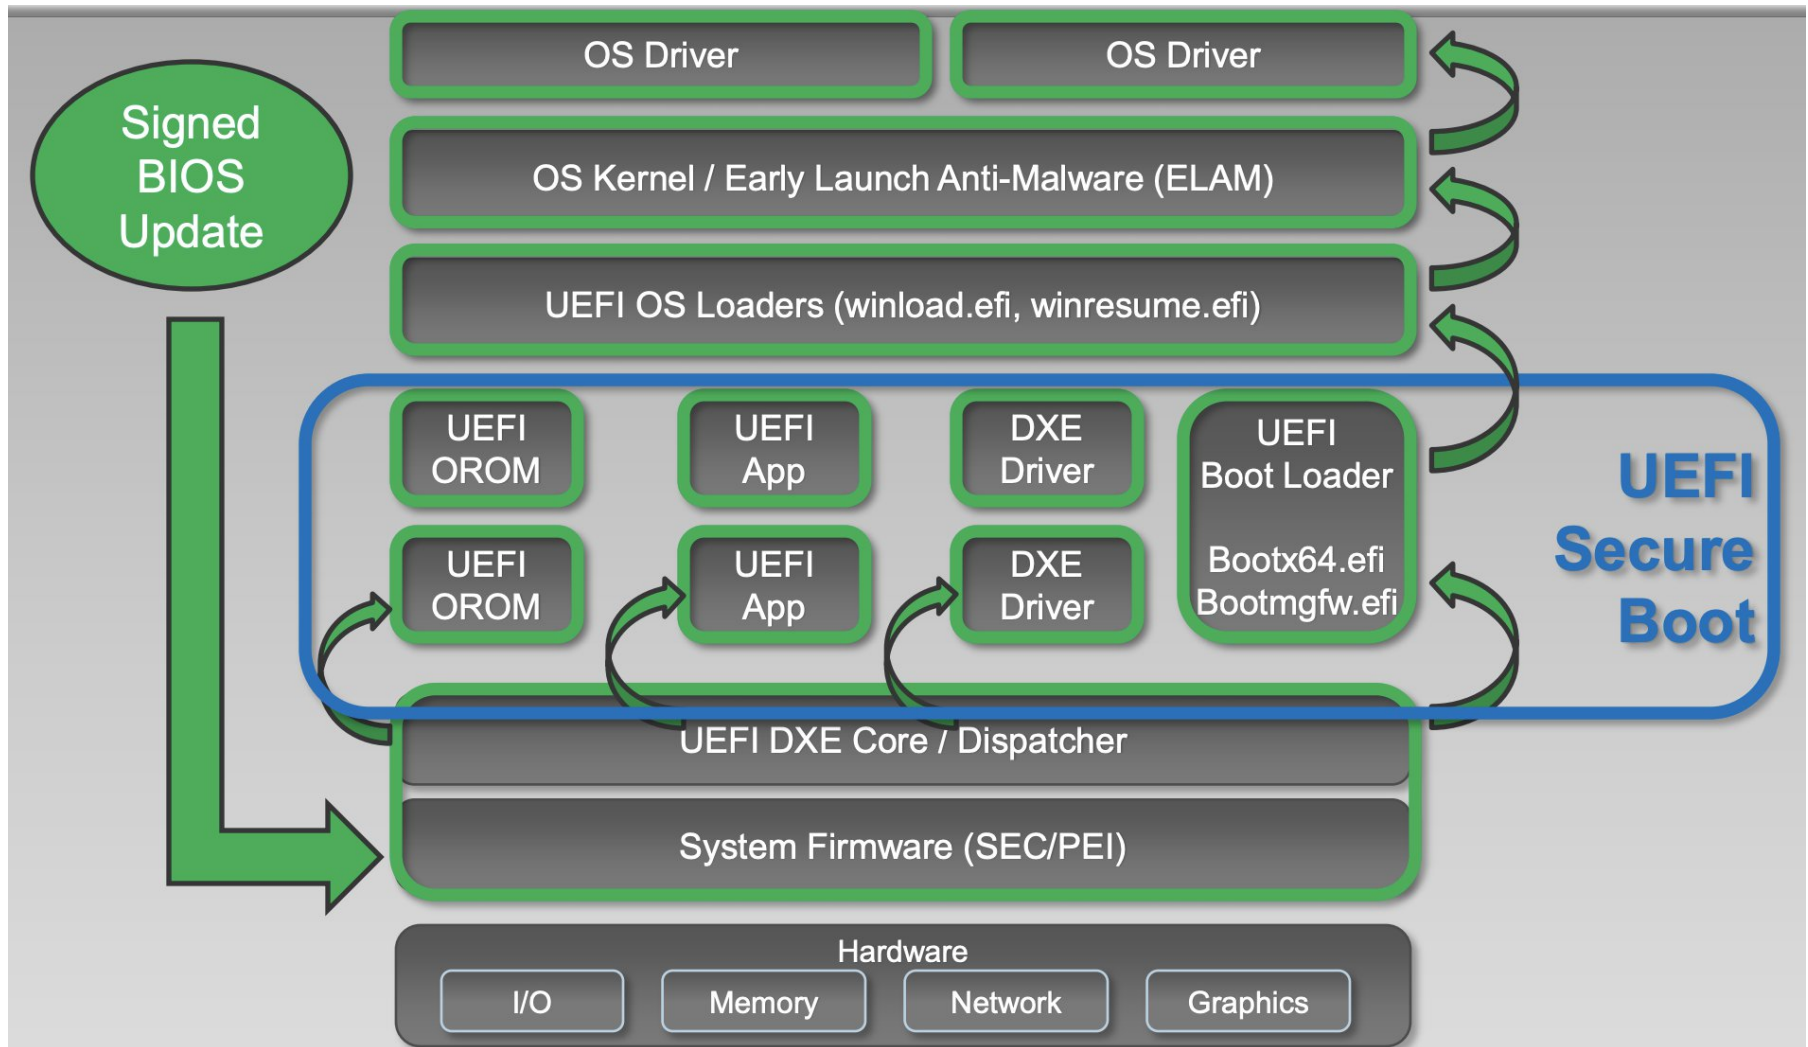
- Control is transferred to TSL phase code
 - TSL = Transient System Load
 - This is typically the OS bootloader loaded from the SSD/HDD
- The OS bootloader loads the OS kernel into memory
- It can still use the Boot Services, set up in earlier phases, available via the UEFI Interface
- Before passing control to the OS kernel, it calls ExitBootServices() via the UEFI Interface
 - Memory holding DXE drivers is freed up
 - Only RT drivers (and SMM code) remain resident and can be used by the OS as runtime firmware services available via the Runtime Interface (or via SMIs)



UEFI Secure Boot

- General idea:
 - Verify if an executable (e.g., OS bootloader) is permitted to load and execute during the UEFI boot process
 - Verification is based on checking digital signatures on code (or checking code hashes)
 - » if executable is signed with an authorized key --» allow
 - » if hash of executable is stored in DB of authorized hashes --» allow
- Caveats:
 - Secure Boot is optional; the way it is managed, enabled or disabled is a decision of the platform manufacturer and the system owner
 - Security boils down to managing authorized keys and hash databases
 - Flash based UEFI components (SEC, PEI, DXE Core) are not verified
 - » they are implicitly trusted
 - » the flash image itself is signed, but it is verified only when loaded into the flash during a firmware update --» see UEFI secure firmware update later...

UEFI Secure Boot – Chain of Trust

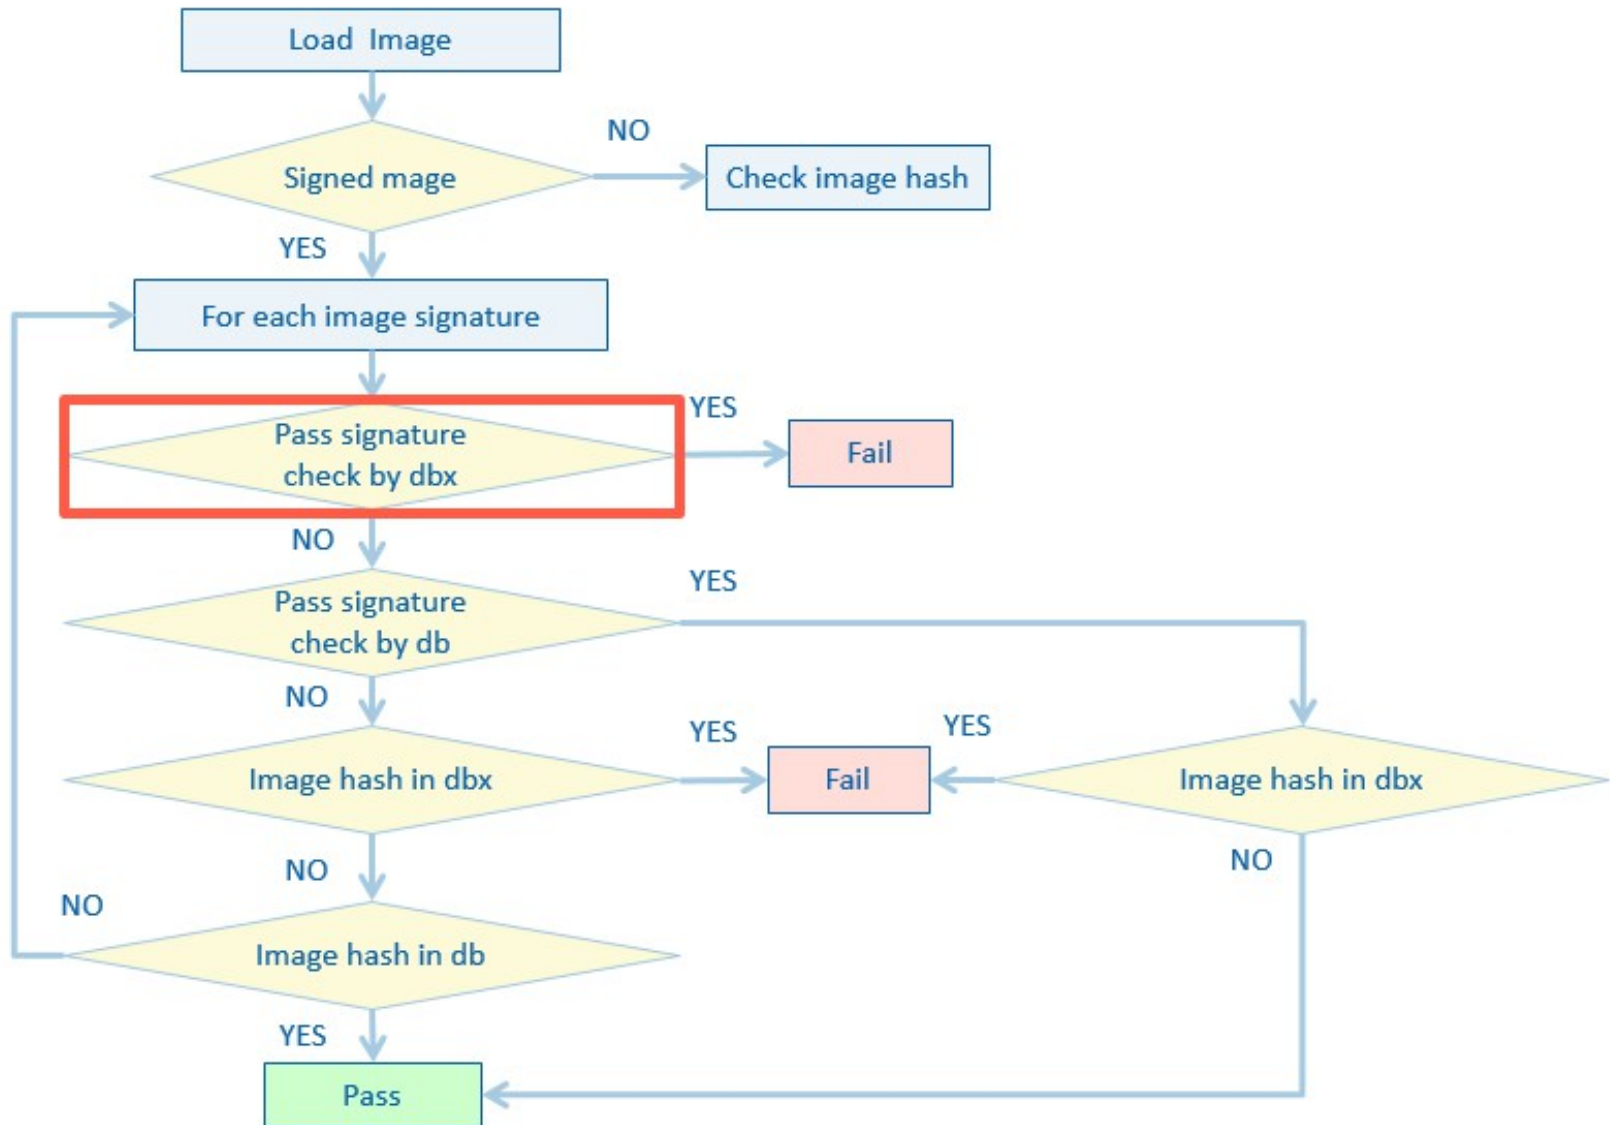


http://c7zero.info/stuff/Windows8SecureBoot_Bulygin-Furtak-Bazhniuk_BHUSA2013.pdf

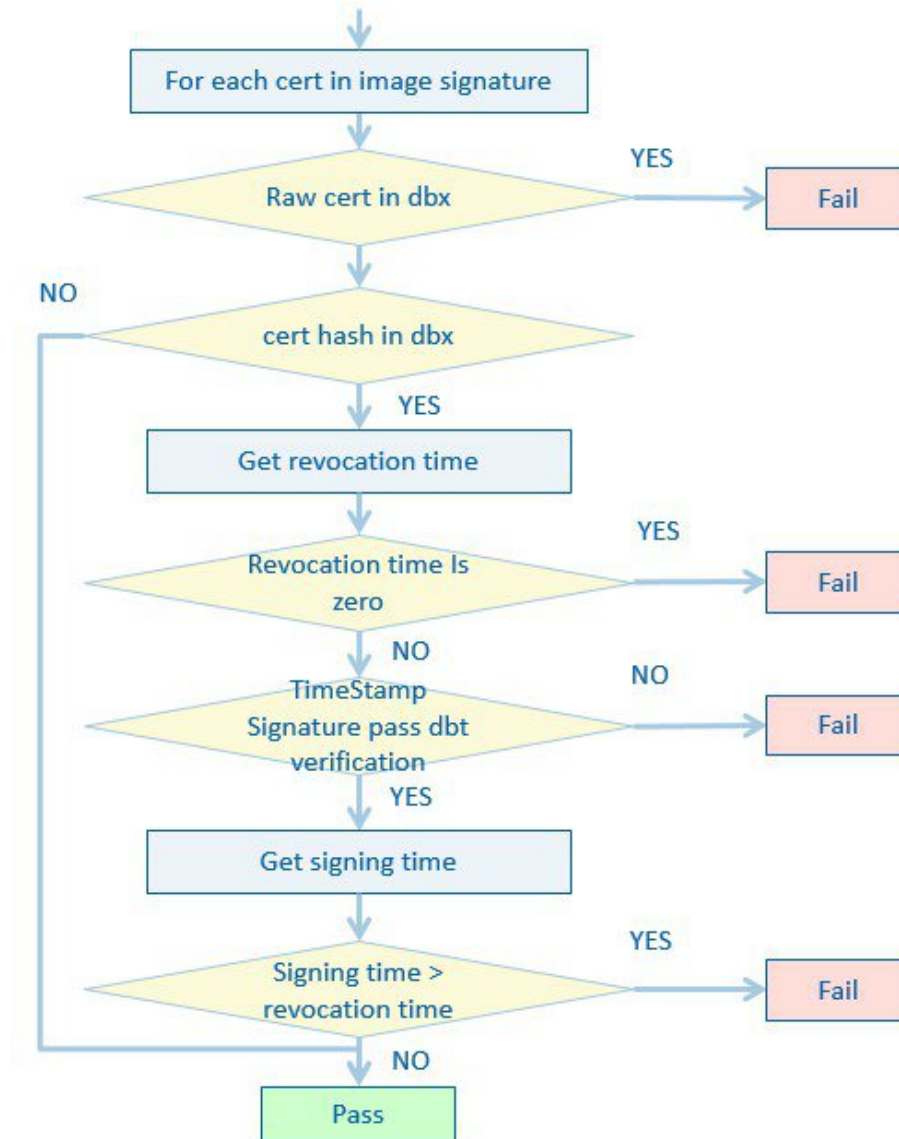
UEFI Secure Boot – Verification process

- Signature and hash checks
 - db file: trusted certificates and hashes
 - dbx file: revoked certificates and untrusted hashes
- The db files are stored in the SPI flash and later copied to memory
- Instead of adding a lot a signer certificates to the db, developers can ask Microsoft a certificate signed by their CA, which can be added to the db
- Linux distributions were also signed by developers' keys signed by the Microsoft CA

UEFI Secure Boot – Verification process



UEFI Secure Boot – Verification process



Secure Firmware Update

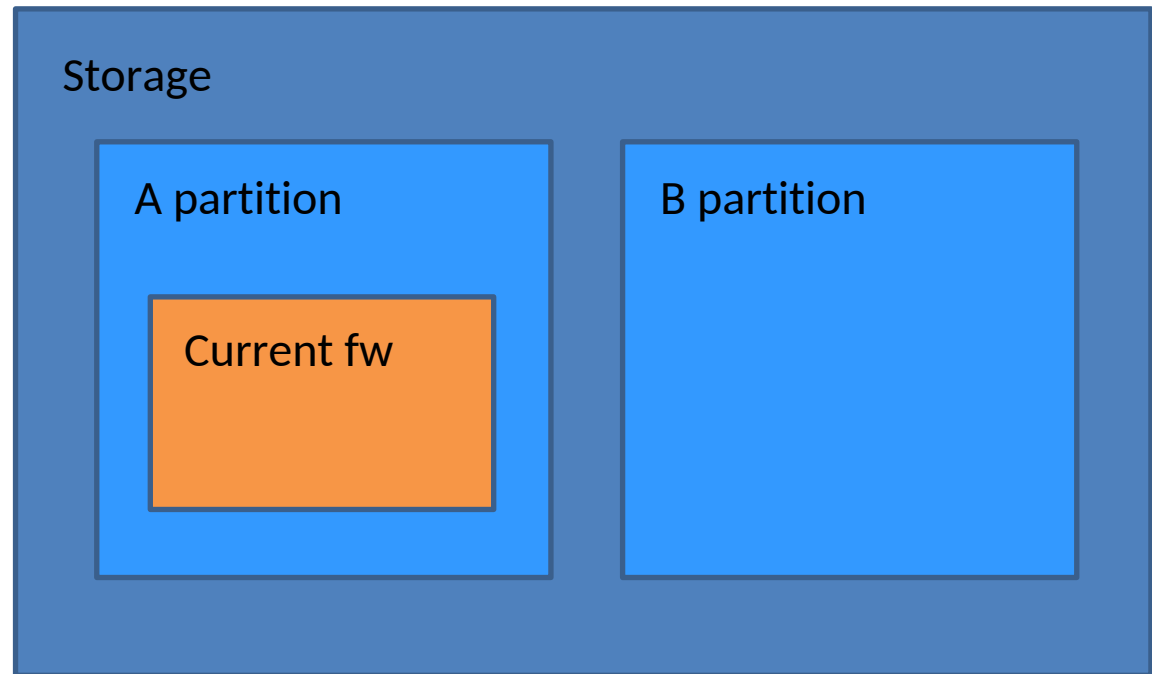
Secure Firmware Update

- Software contains bugs & vulnerabilities
 - Devs fix bugs, new version must be deployed somehow
 - Usual update processes work on a package granularity
 - » Some software components cannot be bundled to packages
- Downloaded components must be integrity-checked
 - Just like in the case of secure boot
- The update mechanism must be fail-safe
 - The system must be able to recover in case of faults
 - Rollback must be possible – use the previous version, if the new cannot be used for some reason
- Rollback protection needed
 - An attacker must not be able to install an older version
 - » Version numbers must be checked & they must be integrity protected as well

A/B partitioning

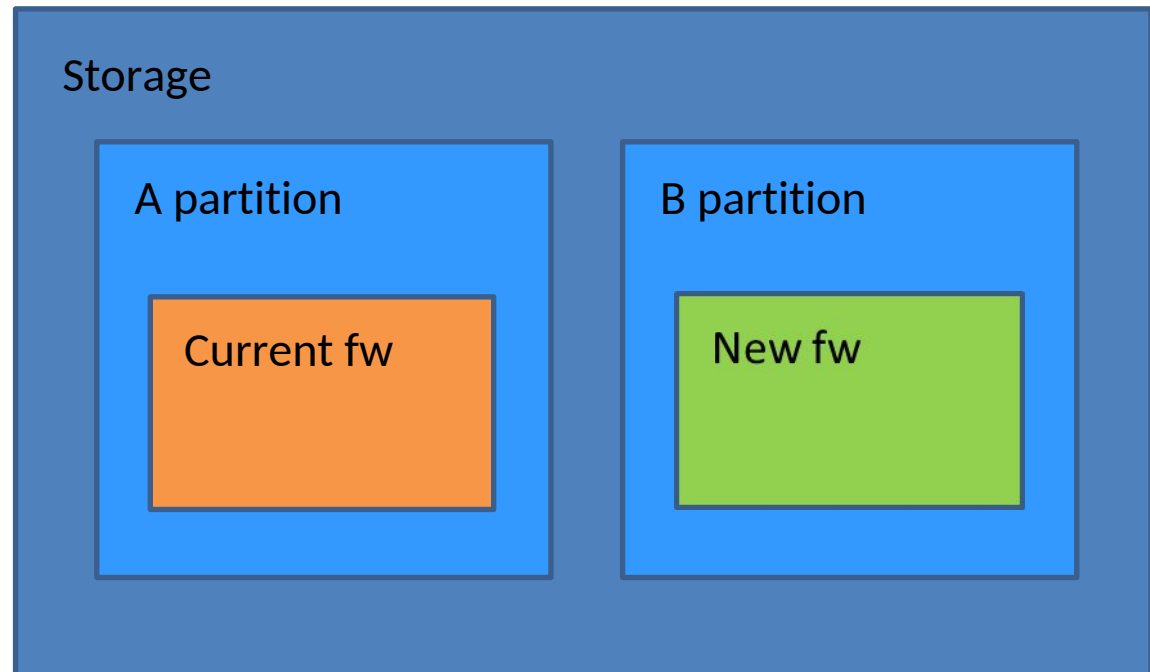
A/B partitioning

- Current firmware downloads new version to free partition



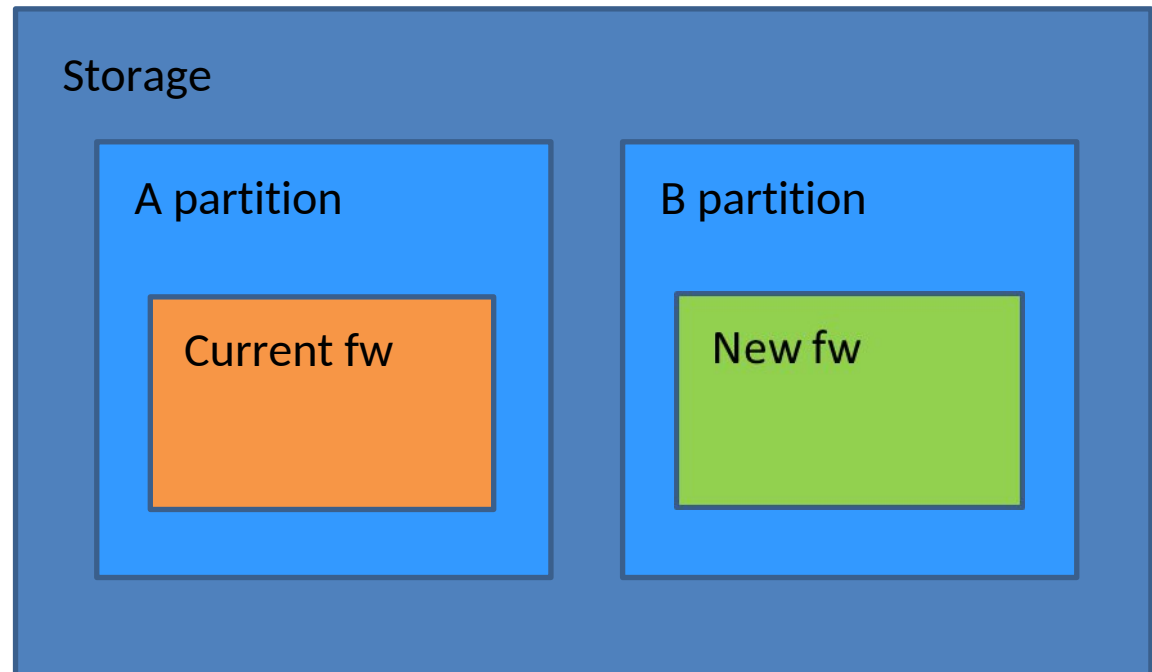
A/B partitioning

- Current firmware downloads new version to free partition



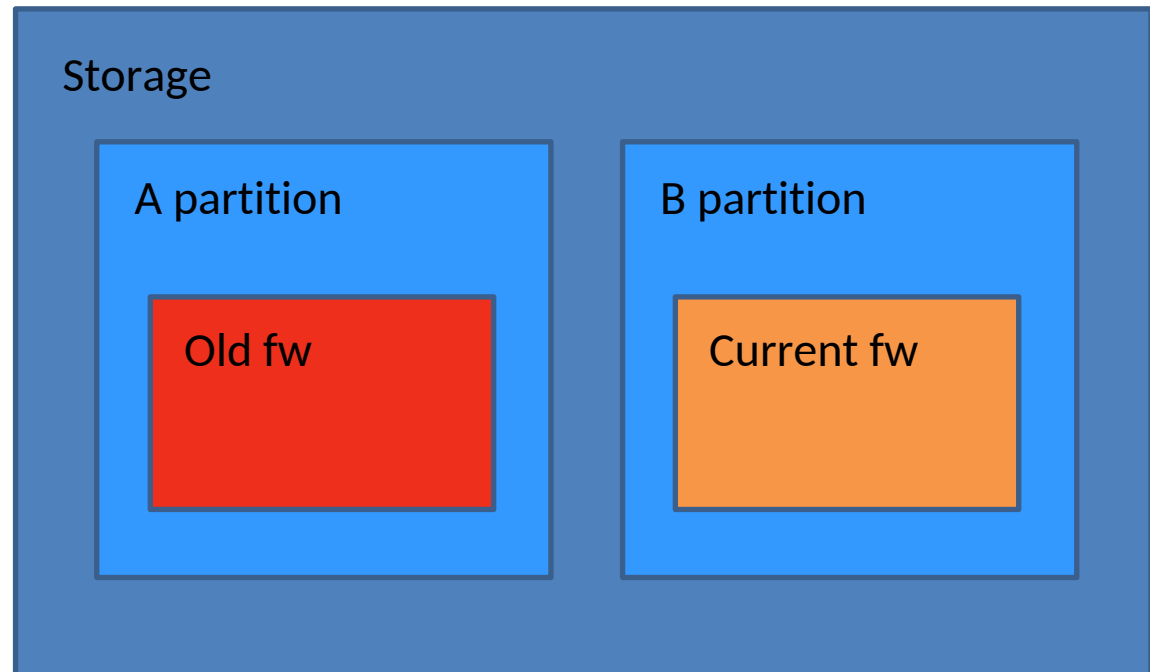
A/B partitioning

- Current firmware downloads new version to free partition
- Device reset, boot into new firmware



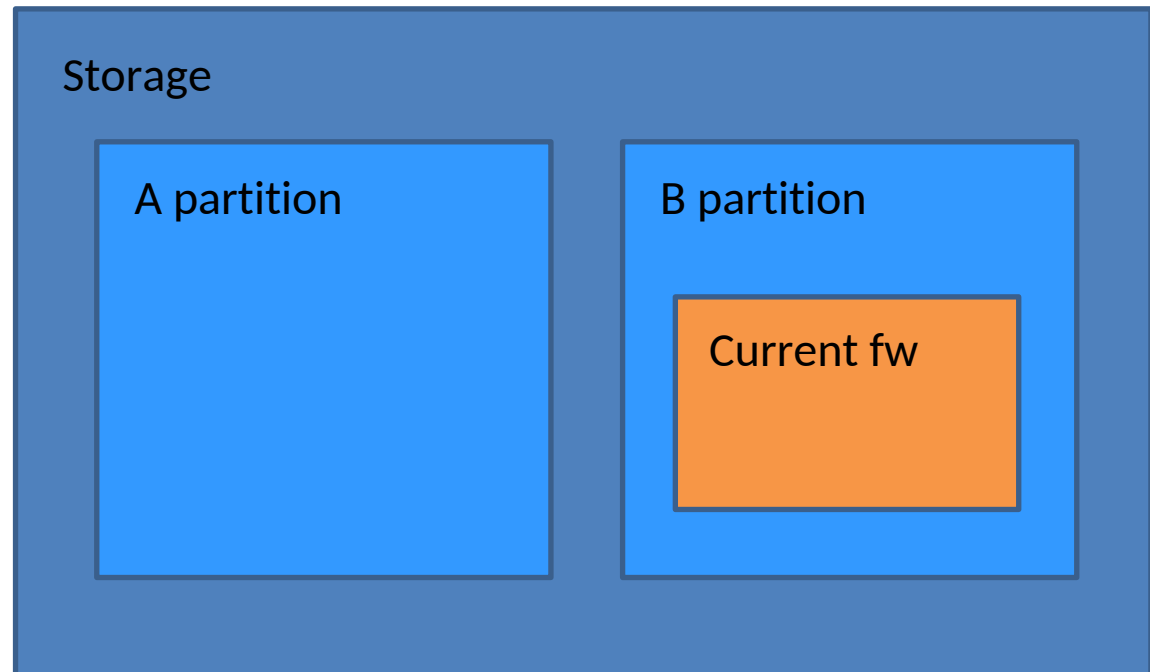
A/B partitioning

- Current firmware downloads new version to free partition
- Device reset, boot into new firmware
 - If successful, old firmware can be deleted



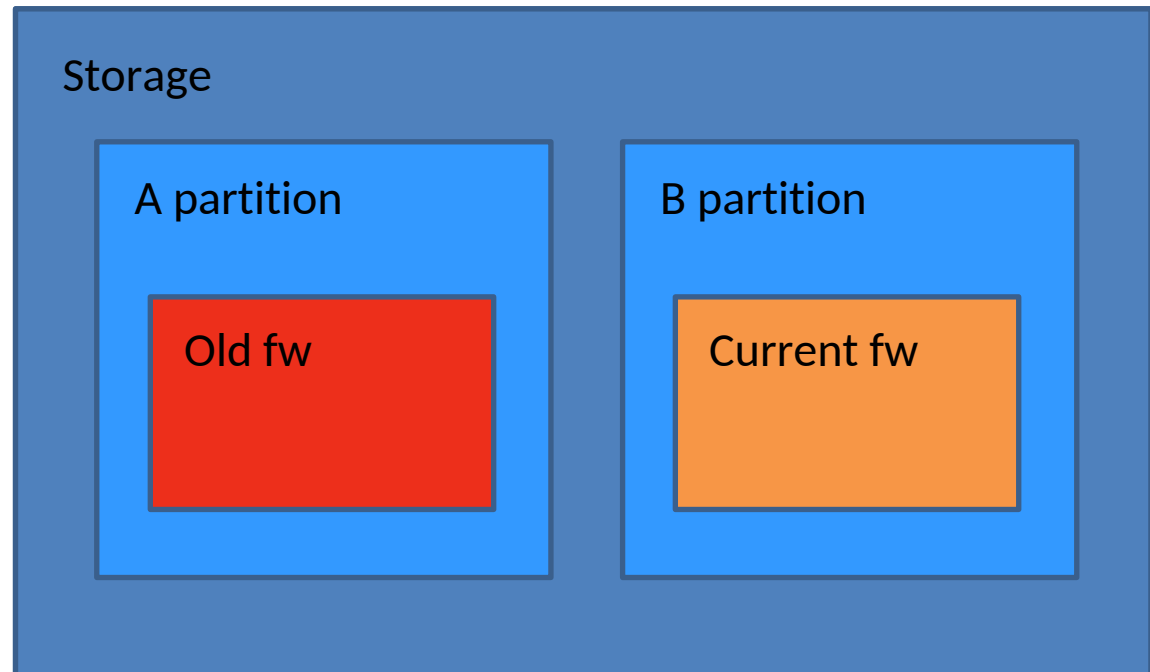
A/B partitioning

- Current firmware downloads new version to free partition
- Device reset, boot into new firmware
 - If successful, old firmware can be deleted



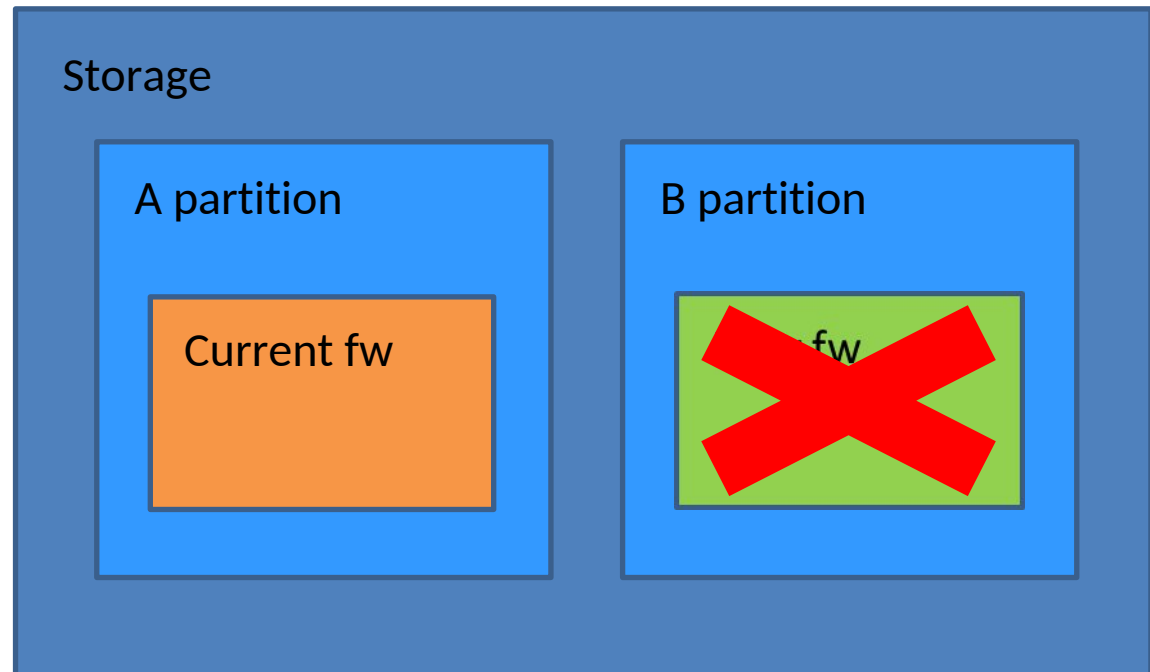
A/B partitioning

- Current firmware downloads new version to free partition
- Device reset, boot into new firmware
 - If successful, old firmware can be deleted
 - If not, rollback to original



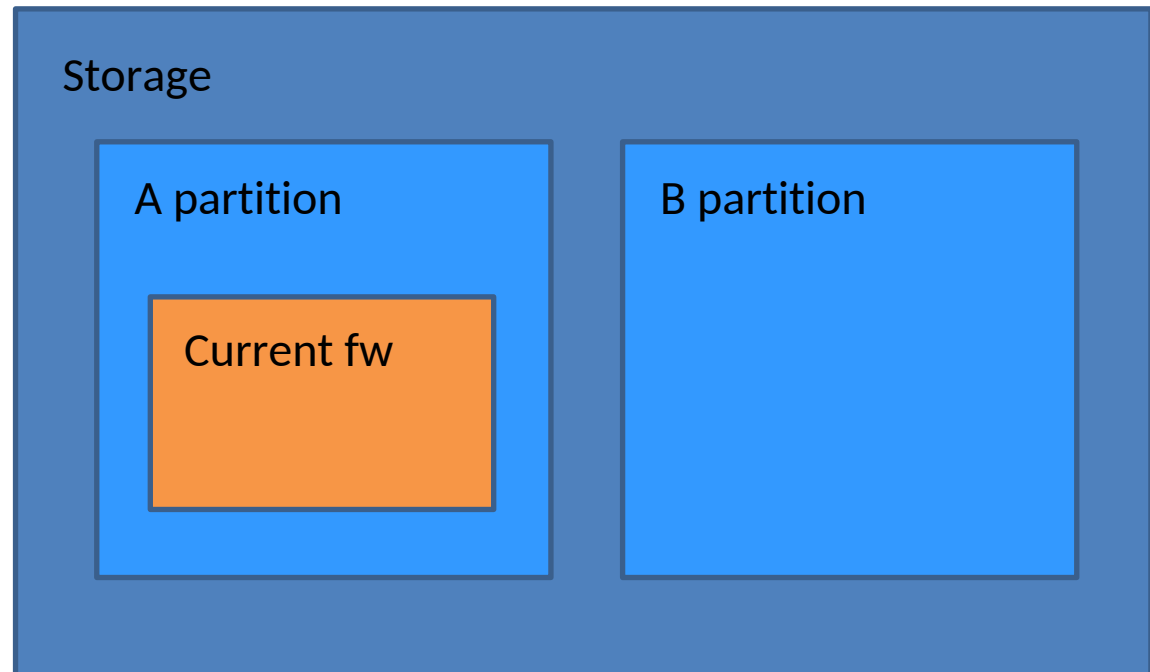
A/B partitioning

- Current firmware downloads new version to free partition
- Device reset, boot into new firmware
 - If successful, old firmware can be deleted
 - If not, rollback to original



A/B partitioning

- Current firmware downloads new version to free partition
- Device reset, boot into new firmware
 - If successful, old firmware can be deleted
 - If not, rollback to original



ESP32 Secure Firmware Update

ESP32 Over The Air (OTA) update

- The device has a partition table that describes the content of the flash
 - Required partitions for OTA update
 - » ota_0, ota_1, to store downloaded images
 - » otadata , to store data about the update process (e.g. which partition to use on next boot)
- Some functions are implemented, but the developer must perform:
 - Checking periodically for new images
 - Download & write the new firmware
 - » Preferably by using a secure channel
 - Check the integrity of the downloaded image
 - » Supported: SHA256 hash
 - Perform self test & cancel rollback in the new image

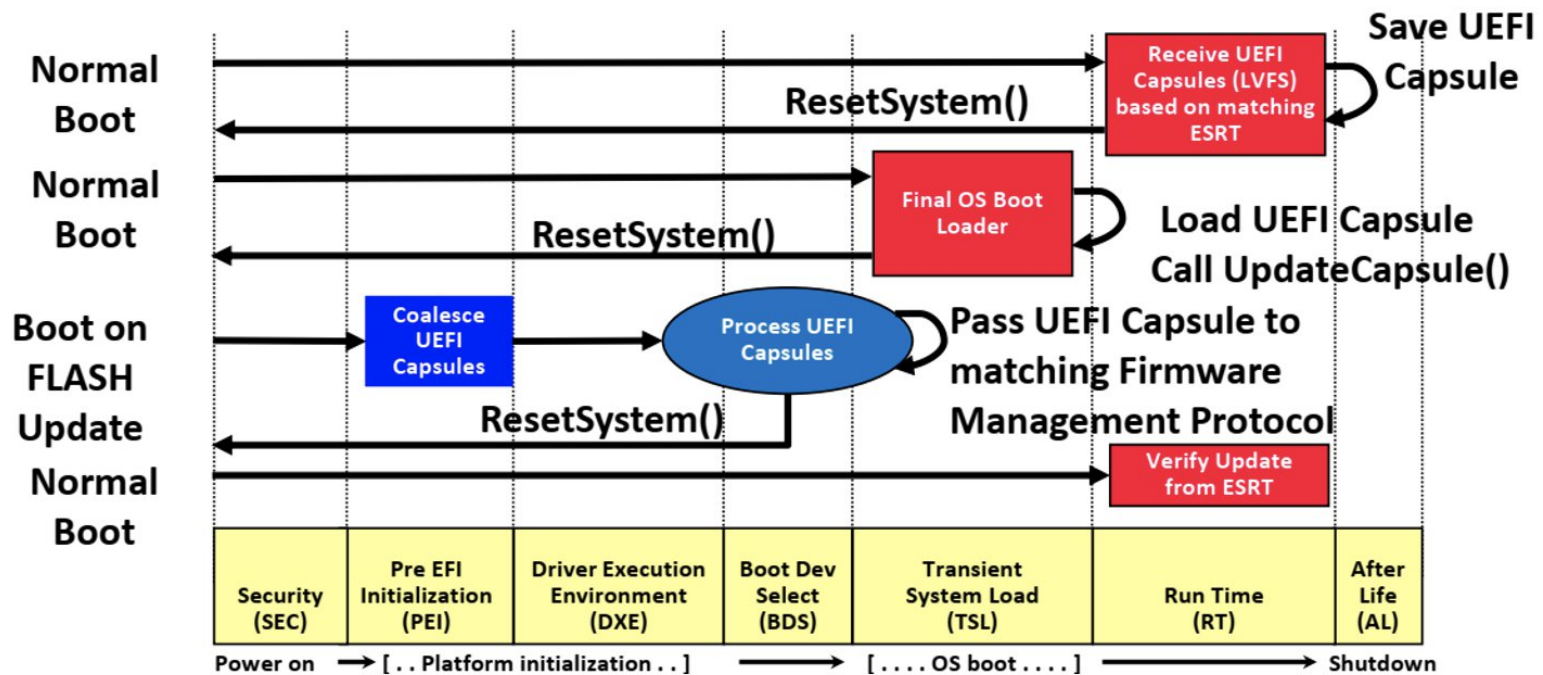
ESP32 Rollback Protection

- All firmware images include a secure version number
 - Must be incremented by the developer before release
 - As it's part of the image, it cannot be changed without changing the hash
- When the partition containing the new image is marked to be executed on next boot, this version is checked
 - If smaller than the current, the partition is erased
- The bootloader checks the version before passing execution to the firmware
 - The partition is not selected for execution, unless the version is greater than or equal to the current version

UEFI Secure Firmware Update

UEFI Secure Capsule Update

- Capsule: a header + data, signed by a key known to the current firmware
- It must be downloaded and copied to a partition accessible to UEFI



UEFI Secure Capsule Update

- HW protection
 - Only the firmware itself can write the SPI flash
- Recovery
 - During update, it is possible to save the content of the SPI flash to disk, to support automatic recovery, in case the freshly downloaded firmware version fails to boot properly
- Rollback protection
 - UEFI firmware contains a table for configuration-related information
 - » EFI System Resource Table (ESRT)
 - The ESRT contains a Lowest Supported Firmware Version field
 - In case of a failure, rollback is possible, but not beyond the Lowest Supported Firmware Version
 - » Rollback to versions with known vulnerabilities can be restricted

Key takeaways

- Secure boot
 - Requires hardware support
 - Each component is checked by its predecessor
 - For the first steps, check what the vendor recommends

- Secure Firmware update
 - Needs to be fail safe
 - » Self test needed
 - » The previous version must be available
 - Must be protected from rollback attacks
 - » Version numbers must be checked & protected

Control questions

- What are the goals of a platform firmware?
- Why is firmware security important?
- What are the phases of the UEFI boot process?
- How is the Root of Trust provided for ESP32?
- What does UEFI stand for? What is standardized by the UEFI Forum?
- How are images verified by UEFI Secure Boot before loading and executing them?
- What keys and databases are used by UEFI Secure Boot? How are they managed?
- What is the goal & the main idea behind A/B partitioning?
- Why self testing is important?
- How can we prevent version rollback attacks?
- What is UEFI Secure Capsule Update?