



Logging and network forensics

Tamás Holczer

CrySyS Lab

Budapest University of Technology and Economics

holczer@crysys.hu

Outline

- Logging (partly based on the presentation of Péter Höltzl)
- Network forensics (partly based on the presentations of Levente Buttyán and András Gazdag)
- ELK stack partly based on the presentation of Gergő Ládi

What is a log

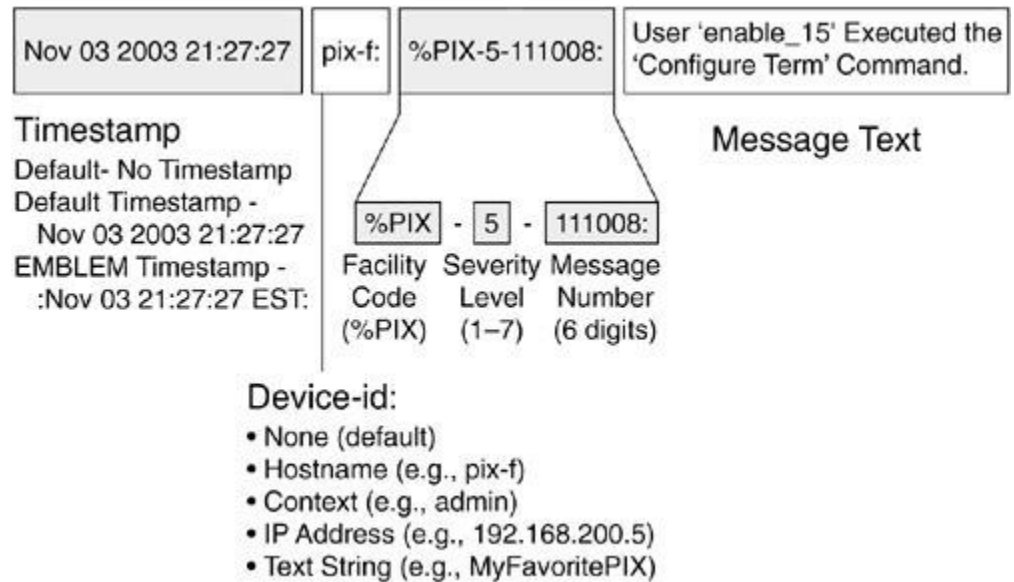
- LOG = record of events (entries)
- Goal of logging:
 - Debugging
 - Performance optimization
 - Authorized and unauthorized activity recording
 - Record of compliance
 - Policy

Who creates logs

- Applications
 - Web server
 - Email server
 - VPN
 - DHCP Server
 - AV
 - ...
- Network devices
 - Firewall
 - Switch
 - ...
- OS
- IDS, IPS
- ...

BSD syslog (RFC3164)

- Developed in 80s
- Orig: part of sendmail
- RFC3164 2001
 - Documents the status
- RFC5424 2009
 - Standardizes syslog
 - Obsoletes 3164



- Device ID: usually hostname (no FQDN)
- Facility: kern, user mail, daemon, auth...
- Severity: 0-Emergency, 7-Debug
- MSG: Latin-1 free text

Problems with Syslog

- UDP 514
- No unique identifier for events
- No acknowledgement
- No security (integrity protection or encryption)
- Timestamp: no year or timezone in many cases
- No multiline messages
- No L7 acknowledgement

- Best effort service no reliability

Syslog API and syslogd

- Applications normally uses the Syslog API
- Syslog events goes to /dev/log
- syslogd collects records from /dev/log and stores them (default: /var/log/syslog) according to a configuration

```
#include <syslog.h>
```

```
syslog (LOG_MAKEPRI(LOG_LOCAL1, LOG_ERROR), "Unable to make  
network connection to %s. Error=%m", host);
```

```
import syslog
```

```
syslog.syslog(syslog.LOG_ERR, "Some error happened")
```

Reliable Delivery for syslog (RFC 3195 2001)

- Based on original RFC3164
- Uses TCP (acknowledgement)
- Cryptographic protection
 - Encryption: TLS_RSA_WITH_3DES_EDE_CBC_SHA
 - Authentication: based on MD5
- Raw profile: single line
- Cooked profile: multiline, xml
- Error codes from HTTP:
 - 200 Success
 - 500 General syntax error
 - ...

New standard for syslog

- RFC5424 2009
- Obsoletes old RFC3164
- RFC5425 - RFC5424 over TLS
- RFC5426 - RFC5424 over UDP
- RFC5427 - PRI definitions
- RFC5848 – digitally signed RFC5424 (SDATA field)
- Well defined timestamp format
- Multiline
- TCP and TLS
- UTF-8
 - <165>1 2003-08-24T05:14:15.000003-07:00 192.0.2.1 myproc 8710
 - - %% It's time to make the do-nuts.

Problems with 5424

- No L7 acknowledgement
- No authentication (only implicit with optional TLS)
- Only optional integrity protection
- Not widely implemented (but: syslog-ng)

Other logging solutions

- Microsoft eventlog
 - EVT API -> file (%SystemRoot%\System32\winevt\Logs*.evtx)
 - Event Viewer
 - Local log facility
 - Remote log: RPC
- SQL (INSERT...)
- Text files (e.g Python import logging)
- CLF (Common Log Format) standard text log format for web server
 - 127.0.0.1 user-identifier frank [10/Oct/2000:13:55:36 -0700] "GET /apache_pb.gif HTTP/1.0" 200 2326
- SNMP (Simple Network Management Protocol)
 - GET/SET Request
 - Trap
 - SNMP v1-2-2c-3 (<3: cleartext community strings, 3: confidentiality, integrity, auth)
- SDEE (Security Device Event Exchange)
 - Mainly for security events
 - Standard of International Computer Security Association
 - Mainly used by Cisco

Structured logging

- JSON (JavaScript Object Notation):
 - `{ "sender" : "michael" "recipient": { "name" : "michael", "name" : "andrea", "name" : "itay" } subject:"I <3 logs" }`
- WELF (WebTrends Enhanced Log file Format):
 - `pri=123 date=2015-08-17T10:10:10.000+01:00
host=test program=pf pid=123 IN=eth0 OUT=
MAC=00:4a:54:c2:f7:e5:00:08:e5:ff:fd:90:08:00
SRC=1.2.3.4 DST=5.6.7.8 LEN=40 TOS=0x00 PREC=0x00
TTL=49 ID=0 DF PROTO=TCP SPT=51777 DPT=80
WINDOW=0 RES=0x00 RST URGP=0`
- XML

Common problems

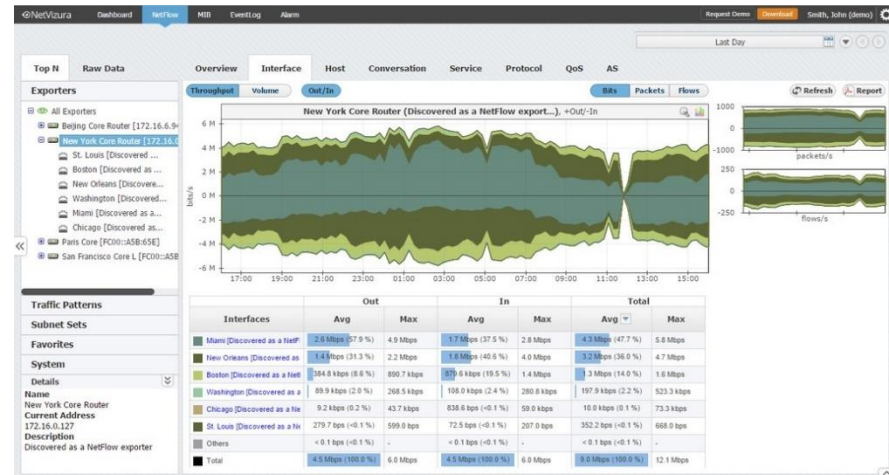
- Different formats
- Not normalized (e.g. timestamp)
- String instead of structured text
- Volume problems
 - High EPS (event per second)
 - Lof of concurrent connections
 - 1 event creates lot of messages

Storage questions

- Local storage
 - No traffic
 - Hard to use
- Central storage
 - Network usage
 - „Safe” place (attacker cannot erase after compromise)
- Mixed storage
 - Locally interesting
 - Locally interesting but without storage (router, switch)
 - Globally interesting
- Encrypted storage?
- Digitally signed storage?

Packet capture

- Full packet capture (mainly for forensics, later in this lecture)
- Flow collection
 - Who communicates with whom at when
 - No payload collected
 - NetFlow / IPFIX (NetFlow v10)
 - Source IP, Destination IP, Source port, Destination port, Time, Header fields...
 - Sender: router, switch, firewall, server...
 - Destination: flow collector
 - Analyser: dashboard, report, alert



The ELK Stack



- ELK = elasticsearch, logstash, kibana
- One of the most popular log management and analytics solutions
- All open source software

The ELK Stack

- Elasticsearch
 - A search and analytics engine, based on Apache Lucene
 - A NoSQL database
 - Has a REST API
 - Sharding and replica support
 - Plugins for analysis, alerting, indexing, etc.
 - Written in Java

The screenshot displays the Google Cloud Platform console for the ELK stack configuration. The 'First cluster' page is active, showing the 'Data' tab with settings for 'gcp.data.highio.1'. The 'Kibana' configuration page is also visible, showing settings for 'gcp.kibana.1'. A 'Summary' panel on the right provides an overview of the cluster's resources and costs.

Summary	
Name	First Cluster
Version	v7.0.1
ES data memory	24 GB
ES data storage	1.25 TB
Total memory	25.5 GB
Total storage	1.25 TB
Hourly rate	\$0.8281
Monthly rate	\$604.51

Architecture

Zone 1

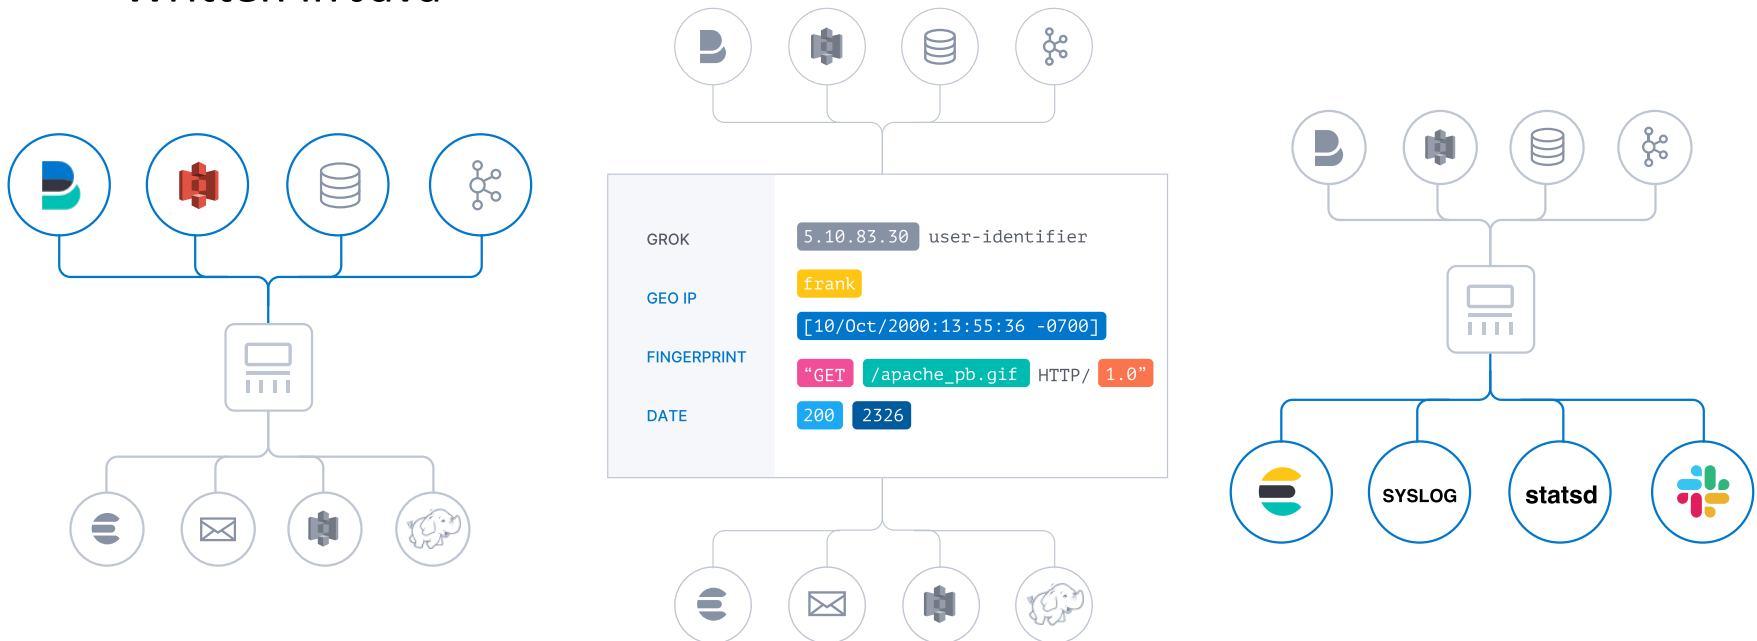
- gcp.data.hl... 8 GB RAM
- gcp.data.hl... 4 GB RAM
- gcp.kibana.1 1 GB RAM
- gcp.apm.1 512 MB RAM

Zone 2

- gcp.data.hl... 8 GB RAM
- gcp.data.hl... 4 GB RAM

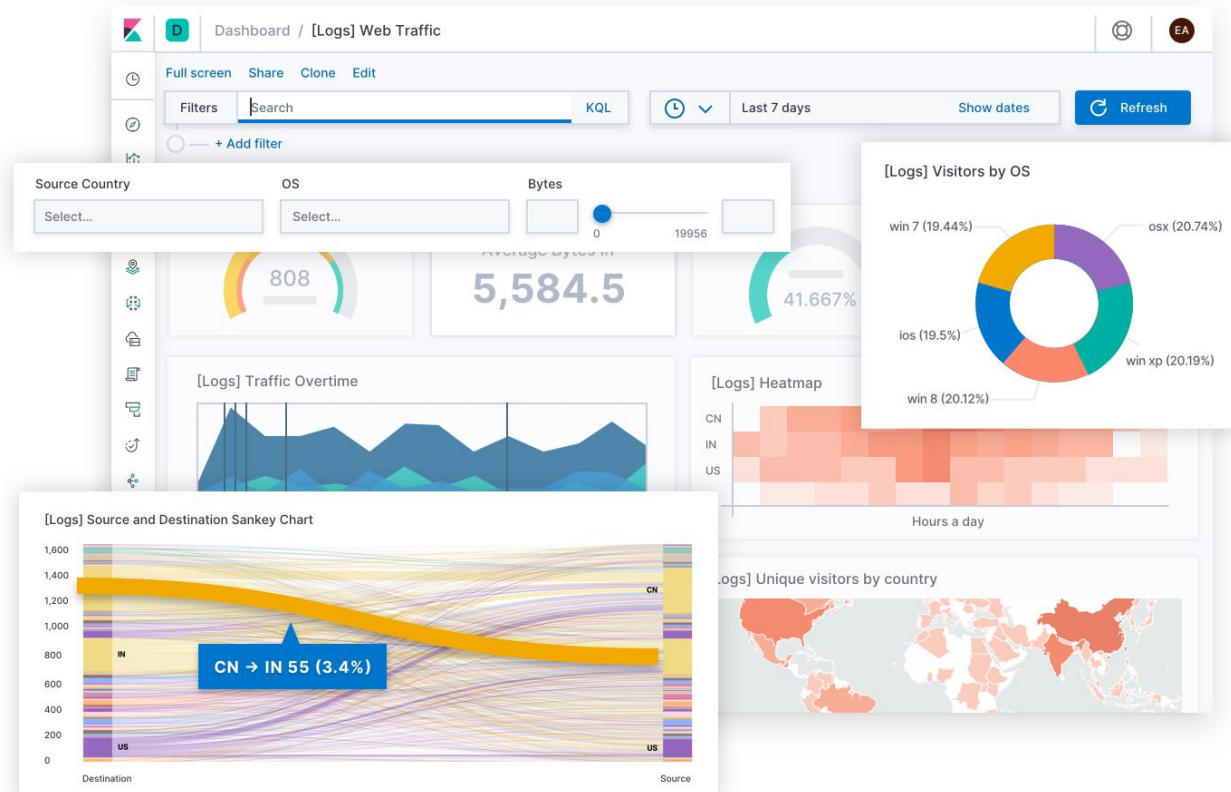
The ELK Stack

- Logstash (Alternative: Fluentd/td-agent)
 - Collects and processes logs from different sources
 - Supports more than 50 different source formats
 - Supports several output formats
 - In our case, it feeds data to Elasticsearch
 - Written in Java



The ELK Stack

- Kibana
 - A web interface to query data in Elasticsearch
 - Data visualization, dashboards
 - Written in node.js



The ELK Stack

- Beats (Alternative: Fluentbit)
 - Originally not part of the ELK stack
 - Collects and feeds extra information (not necessarily logs)
 - Written in Go



Filebeat

Log Files



Metricbeat

Metrics



Packetbeat

Network Data



Winlogbeat

Windows Event Logs



Auditbeat

Audit Data



Heartbeat


Uptime Monitoring



Functionbeat

Serverless Shipper

Logstash vs Beats vs Fluentd vs Fluentbit

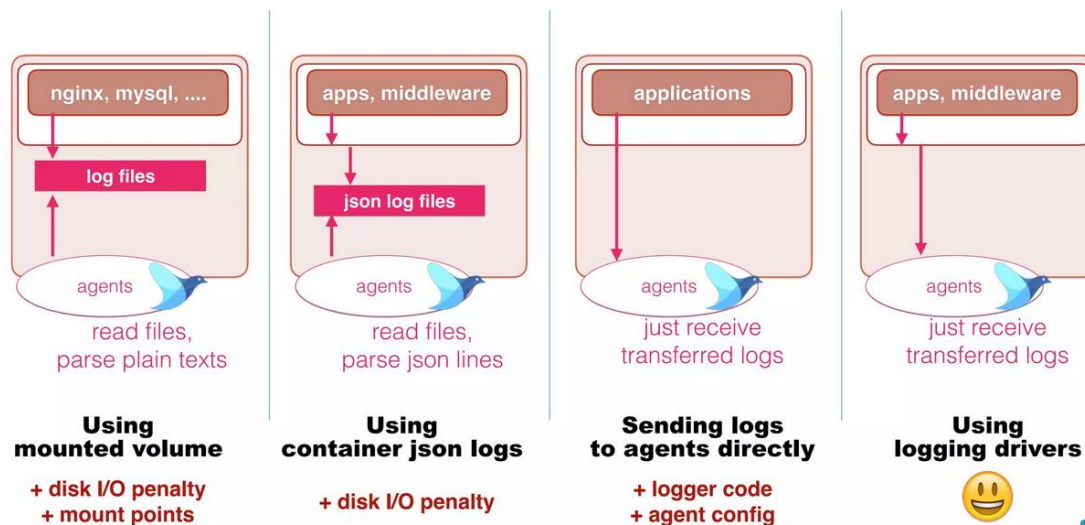


Source: velebit.ai

Distributed logging

- Microservices (containers)
- Everchanging infra (no fixed storage/network/roles)
 - Transfer logs asap
 - Push logs (instead of pull)
 - Inject names/tags into log records (filter logs based on tags later)

How to Ship Logs from Docker Containers



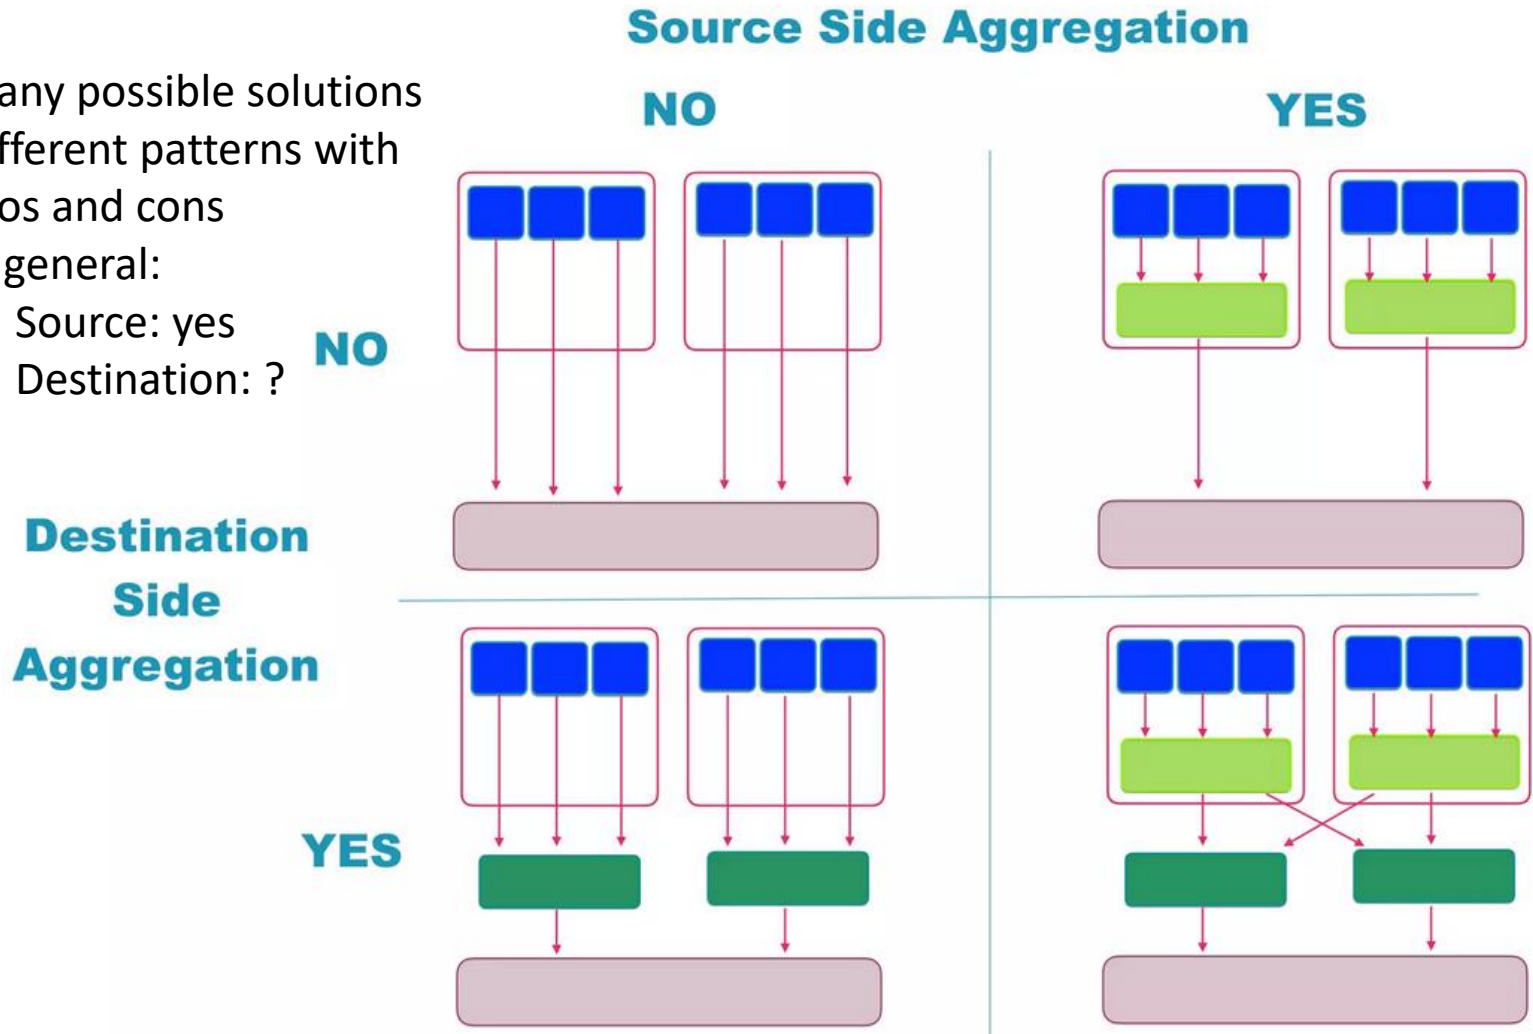
Source:
The Patterns of
Distributed Logging
and Containers
(by Satoshi Tagomori)

Where to aggregate?

Many possible solutions
Different patterns with
pros and cons

In general:

- Source: yes
- Destination: ?



Source: The Patterns of Distributed Logging
and Containers (by Satoshi Tagomori)

Network forensics: Introduction

- network-based evidence can be used to
 - confirm or dispel suspicions surrounding an alleged computer security incident
 - accumulate additional evidence and information
 - verify the scope of a compromise
 - identify additional parties involved
 - determine a timeline of events occurring on the network
 - ensure compliance with a desired activity
- collecting network-based evidence includes
 - setting up a computer system to perform network monitoring
 - deploying the network monitor
 - evaluating the effectiveness of the network monitor
- types of network monitoring:
 - event monitoring
 - » events are alerts that something has occurred on the network
 - trap-and-trace monitoring
 - » records session data summarizing the network activity; includes header fields and flags
 - full content capture
 - » acquiring raw packets collected from the “wire”

Setting up a network monitoring system

- determine your goals for performing the network surveillance
- ensure that you have the proper legal standing to perform the monitoring activity
- acquire and implement the proper hardware and software
- ensure the security of the monitoring platform, both electronically and physically
- ensure the appropriate placement of the monitor on the network
- evaluate your network monitor performance

Determining goals

- goals influence the selection of hardware, software, and filters to be used for collecting network-based evidence, as well as their placement in the network topology
- examples for potential goals:
 - watch traffic to and from a specific host
 - monitor traffic to and from a specific network
 - monitor a specific person's actions
 - verify intrusion attempts
 - look for specific attack signatures
 - focus on the use of a specific protocol
- make sure that the policies in place support these goals
 - e.g., are you allowed to monitor the activity of your employees?

Choosing appropriate hardware

- commercially available network diagnostic and troubleshooting hardware
 - can capture data reliably and at the full rate, but ...
 - lack remote management capabilities and proper storage space
 - usually cost a lot of money
- intrusion detection systems
 - easy to deploy
 - have remote management capabilities and storage space
 - but they cannot perform packet capture reliably
- homegrown solutions
 - easy to customize to fit your needs
 - choose CPU power, amount of RAM, disk space carefully
 - the key issue is to ensure your system has the horsepower required to perform its monitoring function

Choosing appropriate software

- many free tools capture network traffic as well as, or better than, their commercial counterparts
 - e.g., tcpdump, WireShark, Arkime (Moloch)
- selection of the appropriate tool may depend on:
 - which host operating system will you use?
 - do you want to permit remote access to your monitor or access your monitor only at the console?
 - do you want to implement a “silent” network sniffer?
 - do you need portability of the capture files?
 - what are the technical skills of those responsible for the monitor?
 - how much data traverses the network?



Operating system


- choose an operating system with the following features:
 - robust TCP/IP networking stack
 - secure remote access (e.g., via SSH)
 - simple mechanisms for disabling unnecessary services and implementing a local firewall
 - ability to run on many types of hardware, with minimal memory and processor requirements
 - low cost (ideally free)
- Unix/Linux variants are typically good choices
 - e.g., FreeBSD satisfies the above requirements and it is optimized for performance

Stealthy operation

- a ***silent sniffer*** is a system that will not respond to any packets it receives (directed IP datagrams, broadcast, or multicast)
- steps for achieving stealthiness:
 - configure the monitoring interface to speak only TCP/IP
 - » some other protocols, such as NetBIOS, create a lot of traffic that would compromise the location of your monitor
 - » on Windows systems, you need to make sure that you unbind all protocols except for TCP/IP
 - » Unix/Linux systems are generally configured out of the box to communicate with TCP/IP only
 - disable the monitoring interface from responding to ARP packets
 - » most Unix/Linux systems support ifconfig command-line options to turn off ARP on the listening interface
 - » if the monitoring software requires an IP address on the listening interface, assign the system a null IP address (0.0.0.0)
 - alternatively, one can use a one-way Ethernet cable for silent monitoring
 - » disconnect the transmit wires on the network cable
 - » inexpensive, yet very effective
- before deployment of the monitoring system, it is a good idea to run a port scanner (e.g., Nmap) against it, as well as a sniffer detection tool (e.g., Microsoft Promqry or L0pht's AntiSniff)

Deployment of the monitoring platform

- the placement of the network monitor is possibly the most important factor in setting up a surveillance system
 - modern networks are often switched
 - » each frame is forwarded only via the port that connects the segment where the intended destination resides
 - » if the monitor is on another segment, it will not be able to sniff the frame
 - connect the monitor to the SPAN port of the switch
 - » special port via which all traffic is forwarded by the switch
 - if the SPAN port is already used, then use an Ethernet tap
 - » a tap has (at least) three ports: an A port, a B port, and a monitor port
 - » it passes all traffic between ports A and B unimpeded, but it also copies that same traffic to its monitor port



Evaluating the monitoring performance

- things to verify after deployment
 - the traffic monitoring program is executing appropriately
 - the monitoring platform can handle the load
 - the disk isn't filling up rapidly
- e.g., on Unix/Linux, you can use the `df -h` and `top` commands

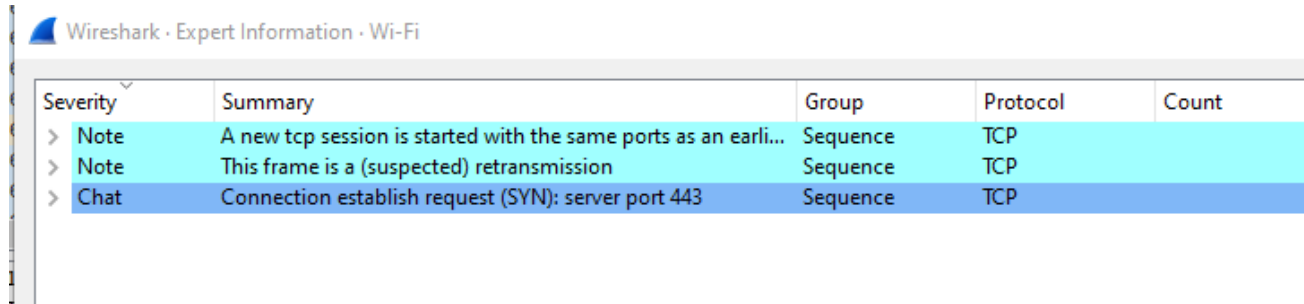


What it can and cannot do

- It cannot be used to map out a network
- It does not generate network data-Passive tool
- Only shows detail information about protocols it understand
- It can only capture data as well as the OS\Interface\Interface driver supports.
- An example of this is capturing data over wireless networks.

Wireshark advanced

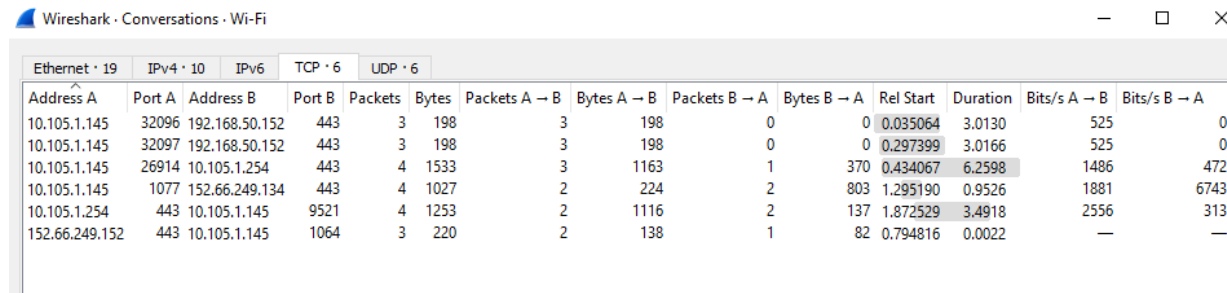
- Analyze / Expert Information



Wireshark · Expert Information · Wi-Fi

Severity	Summary	Group	Protocol	Count
> Note	A new tcp session is started with the same ports as an earli...	Sequence	TCP	
> Note	This frame is a (suspected) retransmission	Sequence	TCP	
> Chat	Connection establish request (SYN): server port 443	Sequence	TCP	

- Statistics



Wireshark · Conversations · Wi-Fi

Ethernet · 19		IPv4 · 10		IPv6		TCP · 6		UDP · 6					
Address A	Port A	Address B	Port B	Packets	Bytes	Packets A → B	Bytes A → B	Packets B → A	Bytes B → A	Rel Start	Duration	Bits/s A → B	Bits/s B → A
10.105.1.145	32096	192.168.50.152	443	3	198	3	198	0	0	0.035064	3.0130	525	0
10.105.1.145	32097	192.168.50.152	443	3	198	3	198	0	0	0.297399	3.0166	525	0
10.105.1.145	26914	10.105.1.254	443	4	1533	3	1163	1	370	0.434067	6.2598	1486	472
10.105.1.145	1077	152.66.249.134	443	4	1027	2	224	2	803	1.295190	0.9526	1881	6743
10.105.1.254	443	10.105.1.145	9521	4	1253	2	1116	2	137	1.872529	3.4918	2556	313
152.66.249.152	443	10.105.1.145	1064	3	220	2	138	1	82	0.794816	0.0022	—	—

- Custom protocol dissector (c/c++ or Lua)
- Analyse TLS based on saved keys (env: SSLKEYLOGFILE)

Summary

- Logging is essential for security and operations
- Many log standards exist
- Logs must be collected and stored in a searchable form
- Full packet capture is a special form of logging

Control questions

- What is the traditional BSD syslog format?
- What are the drawbacks of standard syslog format?
- What extensions are proposed for syslog?
- What is NetFlow/IPFIX used for?
- How to design a full packet solutions?
- What are the parts of the ELK stack? What is their task?
- Distributed logging: <https://www.slideshare.net/tagomoris/the-patterns-of-distributed-logging-and-containers>